

Predicting Formal Protest Petitions Against Housing Development: Evidence from Austin, Texas

M.S. Candidate, Urban Planning
Graduate School of Architecture, Planning and Preservation
Columbia University

May 5, 2026

Abstract

This thesis investigates whether formal protest petitions against proposed housing developments in Austin, Texas can be anticipated before public hearings, using only the administrative and neighborhood records available at the time of application filing. The central question is whether parcel characteristics, neighborhood demographics, and application features observable by planning staff at filing can meaningfully identify which cases are likely to attract organized opposition.

Using discretionary zoning cases from Austin’s pre-HB 24 petition regime (2007–2024), predictive models are trained and evaluated on a temporal holdout that mimics forward administrative use. The results are qualified: filing-date features can rank petition risk above chance—the best model achieves moderate out-of-distribution discrimination with a wide uncertainty interval (PR-AUC = 0.718, 95% CI: [0.48, 0.91])—but the rarity of petitioned cases and the model’s limited external validity mean predictions are best used as relative triage rankings, not as case-level probabilities for high-stakes decisions.

The core finding is that administrative data at filing carries real predictive signal for petition risk, yet that signal is too uncertain and context-dependent to support uncritical case-level decisions. The broader contribution is twofold: first, a methodology—time-stamped features, temporal holdout evaluation, and explainability audits—that is portable to other discretionary review systems; and second, the identification of **neighborhood agency as a predictive bound**. The study finds that the volatility of performance across petition thresholds serves as empirical evidence that mobilization is a strategic social choice that cannot be reduced to mechanistic administrative records.

Acknowledgments

I would like to express my deepest gratitude to my thesis advisor, Dory Kornfeld, for her guidance, support, and expertise throughout this research. Her insights into urban planning and policy were invaluable in shaping this work.

I am profoundly grateful to my thesis readers for their time, expertise, and thoughtful feedback: Matthew Bauer, Adjunct Associate Professor; Josh Begley, Adjunct Assistant Professor; James Piacentini, Adjunct Assistant Professor; Stijn Van Nieuwerburgh, Earle W. Kazis and Benjamin Schore Professor of Real Estate and Co-Director of the Paul Milstein Center for Real Estate at Columbia Business School; and Adam Vosburgh, Associate Director and Adjunct Assistant Professor. Their interdisciplinary perspectives spanning computational design, real estate finance, and urban planning enriched this research immeasurably.

I would like to thank Professor Weiping Wu, Vice Provost for Academic Programs and Director of the Urban Planning Program at Columbia University, for her leadership of the program and support of my academic journey.

I am grateful to the Columbia University Graduate School of Architecture, Planning and Preservation (GSAPP) for providing the academic environment and resources that made this research possible.

Finally, I thank my colleagues, peers, and the broader Columbia community for their intellectual engagement, encouragement, and support throughout this journey.

Contents

1	Introduction	7
2	Literature Review	7
3	Outcome Definition and As-of Information Setup	8
3.1	Analytical Panel Structure	8
3.2	Case Universe and Sample Selection	11
3.3	Primary Outcome Definition	13
3.4	Statutory Context: The Legal Petition Mechanism	13
3.5	Feature Libraries: Policy Forecasting vs. Explanatory Research	13
3.6	Illustrative Panel Structure	15
4	Modeling Strategy	15
4.1	Formal Protest Petition Model	15
4.1.1	Primary Model Selection	16
4.2	Model Evaluation Strategy	16
4.3	Outcome Scope	16
4.4	Feature Attribution and Semantic Clustering	16
5	Filing-Date Primary Results	17
5.1	Predictive Performance Across All Horizons	17
5.2	Predictors of Petition Filing	20
5.3	Attribution Stability and Meta-Attribution	20
6	Institutional Context	29
6.1	Threshold-Based Reduced-Form Discontinuity on Reconstructed Petition Share	29
7	Qualitative Planning Context	30
8	Limitations	32
9	Conclusion	33
	References	34
A	Qualitative Documentation	39
A.1	Interview Recruitment Overview	39
A.2	IRB Protocol Summary	41
A.3	Semi-Structured Interview Guide	42
B	Quantitative Methods and Robustness	43
B.1	Expanded Methods Appendix: Technical Details	43
B.1.1	Hyperparameter Search Grids	43
B.2	Temporal Drift Threshold Sensitivity	45
B.2.1	Threshold: >0%	46
B.2.2	Threshold: >5%	50
B.2.3	Threshold: >10%	54

B.2.4	Threshold: >15%	58
B.2.5	Threshold: >20%	62
B.2.6	Threshold: >25%	66
B.3	SHAP Beeswarm Exhibit	70
C	Institutional Context and Causal Overlays	71
C.1	Institutional Geographic Overlay Analyses	71
D	Policy Event Studies	74
D.1	HOME Initiative Event Study	74

List of Figures

1	Spatial Distribution of Zoning Cases, 2007–2024	10
2	200ft Parcel Buffer Geometries	11
3	Schematic of the Austin Zoning Process	12
4	Sample Formal Protest Petition Record	14
5	Filing-Date Precision-Recall Curves by Model Class	18
6	Temporal Drift Evidence (Target Horizon)	22
7	Temporal Drift Evidence (OOD Offset)	23
8	Comparative Semantic Predictor Reliance between Architectural Families	24
9	Meta-Attribution Behavioral Clustering	26
10	Performance-Weighted Meta-Attribution Clustering	27
11	Attribution Stability: Continuous Divergence Tracing	28
12	Adjacent-Window Attribution Rank Stability	29
13	Threshold-Based Reduced-Form Discontinuity on Reconstructed Petition Share	30
14	Supplementary Hyperparameter Search Results	44
15	Algorithmic Predictability across Petition Thresholds	45
16	SHAP Beeswarm for the Filing-Date Model	70
17	Identification Diagram: Threshold-Based Reduced-Form Discontinuity	71
18	Estimated Descriptive Interaction Effects by District	72
19	Placebo Estimates Across Earlier Election Cycles	73
20	HOME Phase 1 Event-Study Coefficients (Reconstructed Petition Outcome)	74

List of Tables

1	Annual Distribution of Discretionary Zoning Cases, 2007–2024	10
2	Historical Panel Descriptive Statistics	12
3	Sample Selection from Raw Records to Analytic Sample	12
4	Exemplar State Protest Petition Thresholds	15
5	Semantic Feature Aggregation Mapping	17
6	Temporal Predictive Drift (PR-AUC decay)	19
7	Temporal Predictive Drift (PR-AUC lift)	21
8	Temporal Drift (Max-of-Family)	24
9	Archetypal Family Attribution	25
10	Performance-Weighted Archetypal Attribution	28
11	Temporal Predictive Drift (Threshold: >0%)	47
12	Temporal Predictive Drift Lift (Threshold: >0%)	48
13	Temporal Predictive Drift Family Dominance (Threshold: >0%)	49
14	Temporal Predictive Drift (Threshold: >5%)	51
15	Temporal Predictive Drift Lift (Threshold: >5%)	52
16	Temporal Predictive Drift Family Dominance (Threshold: >5%)	53
17	Temporal Predictive Drift (Threshold: >10%)	55
18	Temporal Predictive Drift Lift (Threshold: >10%)	56
19	Temporal Predictive Drift Family Dominance (Threshold: >10%)	57
20	Temporal Predictive Drift (Threshold: >15%)	59
21	Temporal Predictive Drift Lift (Threshold: >15%)	60
22	Temporal Predictive Drift Family Dominance (Threshold: >15%)	61
23	Temporal Predictive Drift (Threshold: >20%)	63
24	Temporal Predictive Drift Lift (Threshold: >20%)	64
25	Temporal Predictive Drift Family Dominance (Threshold: >20%)	65
26	Temporal Predictive Drift (Threshold: >25%)	67
27	Temporal Predictive Drift Lift (Threshold: >25%)	68
28	Temporal Predictive Drift Family Dominance (Threshold: >25%)	69
29	Summary of Quasi-Experimental Estimates	73

1 Introduction

The United States faces a persistent housing affordability crisis, with local barriers to new housing supply drawing increasing scrutiny from policymakers and researchers. While zoning laws are often cited as primary constraints, much of the friction occurs through discretionary political processes, including public hearings, protest petitions, and city council negotiations, that can delay or block individual projects. Understanding how these mechanisms operate in practice is essential for diagnosing the roots of housing scarcity and for designing effective policy interventions. Austin, Texas, provides a compelling case study: rapid population and employment growth between 2010 and 2024 intensified supply constraints, fueling conflict between neighborhood preservation and regional development goals. Austin's contemporary land-use politics also sit atop a longer history of exclusionary planning, racial segregation, and uneven redevelopment stretching from the 1928 city plan into the present [35, 64, 65, 51].

One of the most consequential mechanisms for neighborhood opposition in Austin has been the formal protest petition, enabled under Texas Local Government Code Section 211.006 [58]. Prior to 2025, if property owners representing at least 20% of the land area within 200 feet of a proposed zoning change signed a petition, and the municipal clerk certified its validity, the proposal required a three-fourths supermajority vote in city council to proceed. This statutory process created a powerful tool for delaying or blocking new housing, reinforcing the influence of organized local opposition [21]. In 2025, House Bill 24 replaced this system with a differentiated, multi-tier framework codified under Texas Local Government Code Section 211.0061 [59, 31, 57], but this study focuses on the pre-HB 24 regime, when the protest petition was a central feature of Austin's land use politics.

This thesis asks whether it is possible to predict, at the time of filing, which housing development proposals in Austin will trigger a measured threshold-crossing protest petition, specifically, whether administrative, spatial, and neighborhood features available before public hearings can identify cases likely to face this form of organized opposition. The focus is on forecasting the statutory petition event, not all forms of opposition or the final legal certification by the clerk.

The outcome of interest is a "measured threshold-crossing petition," defined as an event where the reconstructed petition-share calculation exceeds the statutory 20% threshold based on public records available at filing. This measure is distinct from the clerk-certified petition, which is not directly observed. The thesis does not attempt to predict all forms of opposition or the final legal outcome, but rather assesses whether administrative, spatial, and economic features at the time of application can meaningfully signal the risk of this specific institutional event. The main contribution is to clarify the predictive value and limitations of these features for anticipating organized opposition in a high-stakes land use context.

To supplement the quantitative analysis, the study incorporates semi-structured interviews with planning professionals. These interviews provide institutional context, highlighting informal negotiations, strategic petition timing, and neighborhood mobilization, that administrative data alone cannot capture. While not used to validate model performance, these perspectives help situate the findings within the realities of planning practice and underscore the ethical and governance considerations surrounding predictive tools in land use decision-making.

2 Literature Review

This literature review synthesizes three core themes that motivate the thesis's methodological approach.

First, land-use regulation is widely recognized as a binding constraint on housing supply, driving a wedge between construction costs and market prices [25, 24, 28, 49]. Measures of regulatory stringency such as the Wharton Residential Land Use Regulatory Index document substantial cross-metropolitan variation in local land-use controls [29, 27]. Fischel's "homevoter hypothesis" [20] posits that homeowners mobilize politically to protect property values, often leveraging zoning to preserve neighborhood exclusivity;

McCabe [41] extends this logic by linking local politics to household wealth accumulation through homeownership. Empirical work by Einstein, Palmer, and Glick [17] documents "participatory distortion," where older, wealthier, long-tenured homeowners dominate public processes and systematically oppose new housing. Hankinson [32] finds that renters, too, may resist nearby development due to displacement fears, reinforcing the spatially localized nature of opposition. Recent syntheses of NIMBY and YIMBY politics show that these conflicts are embedded in broader coalitions over land use, growth, and neighborhood change rather than reducible to a single homeowner motive [7]. At the same time, supply-side reform arguments coexist with cautions that new supply alone will not fully resolve affordability and displacement pressures in high-cost markets [30, 4].

Second, the power of participatory distortion is magnified in discretionary review systems, where individual projects require site-specific approval rather than proceeding "by right" [39, 33]. Research on housing opposition shows that neighborhood resistance is heterogeneous, often mixing status-quo protection, administrative claims, environmental review, and distributive conflict rather than a single undifferentiated NIMBY logic [48, 53]. In Texas, the protest petition statute empowered adjacent property owners to trigger supermajority council votes, making it a potent institutional veto point [58, 21]. Supermajority requirements systematically favor the status quo [10], amplifying the influence of organized opposition. In Austin, measured threshold-crossing petitions have been especially common for rezonings that increase residential density, highlighting the asymmetric impact of these institutional tools on housing supply.

Third, while much of the urban planning literature focuses on causal inference, estimating the effects of zoning changes or protest petitions on outcomes like land value or project delay, addressing discretionary gridlock also requires anticipating where opposition will arise. Shmueli [55] and Mullainathan and Spiess [44] distinguish between explanatory and predictive modeling, emphasizing that predictive models can inform policy and administrative triage even when causal mechanisms are complex or uncertain. Kleinberg et al. [34] formalize the "predictive policy question," where optimal decisions depend on forecasting rather than explanation. In the context of Austin's protest petition regime, the relevant question is not just why opposition occurs, but whether it can be predicted in advance using only the information available at the time of filing. This predictive framing aligns with emerging best practices in public policy research [1]. It also raises public-sector governance questions about opacity, proxy discrimination, and sociotechnical context that cannot be resolved by predictive performance alone [18, 47, 2, 54].

3 Outcome Definition and As-of Information Setup

The empirical foundation is a time-stamped spatial panel that reconstructs what information was available at each historical prediction point. Every record anchors to an explicit as-of date, ensuring that backtests use only Census releases, appraisal rolls, and case documents published before the prediction date.

3.1 Analytical Panel Structure

The core analytics rely on the filing-date analytic sample, which merges spatial and demographic variables into single case-level observations (with temporally matched covariates):

1. **Zoning Application Records:** Case ID, milestone dates (filing, notice, petition deadline, planning commission hearing, council reading, withdrawal), requested zoning changes, unit limits, target FAR, and disposition status [15].
2. **Spatial Geometries:** Parcel IDs, geometries, acreage, frontage, and 200ft buffers used to aggregate nearby properties (council districts and overlapping geometries) [15].

3. **Parcel Appraisals:** Annually indexed TCAD appraisal variables for parcels used to compute variables within the 200-foot buffer: land value, structure age, homestead exemption rates, and owner-occupancy shares [61].
4. **Neighborhood Demographics:** Tract and block group ACS variables constrained to their historical 5-year release dates: rent burdens, tenure splits, demographic composition, and displacement vulnerability indicators [63].
5. **Institutional Calendar:** Council-term and policy-timing indicators, including the 2022 council turnover, Austin’s 2023 HOME Initiative adoption [14], and the HB 24 effective date (September 1, 2025) [57].

These components are assembled from City of Austin open-data records, Travis Central Appraisal District appraisal files, U.S. Census Bureau ACS 5-year estimates, Travis County certified election results, and statutory materials governing the protest-petition regime [15, 61, 63, 62, 58, 57].

Because municipal databases frequently overwrite preliminary timestamps with final disposition dates, milestone dates are reconstructed through two pathways. First, explicit dates are parsed from scraped historical council agendas and PDF case files. Second, where explicit text is unavailable, dates are imputed by back-calculating from end-dates using minimum statutory constraints mandated by Texas Local Government Code and Austin zoning ordinances (for example, statutory mail-notice minima). This parsed-versus-imputed distinction is carried into interpreting the results, as timing uncertainty introduces measurement error.

Supporting tables log meeting events (staff and commission recommendations), speech comments (speaker stance and segmented text), petition records (validity and signature counts), and vote records (case-by-councilmember dynamics).

Annual Zoning Application Volume (2007-2024)




Figure 1: Spatial distribution of zoning cases across the Austin municipal core ($N = 6783$). *Generated: 2026-04-13*

Table 1: Annual Distribution of Discretionary Zoning Cases, 2007 to 2024

Note: Updated: 2026-04-13

Year	Cases	Year	Cases	Year	Cases
2007	288	2013	366	2019	460
2008	270	2014	390	2020	362
2009	265	2015	454	2021	476
2010	297	2016	422	2022	436
2011	369	2017	483	2023	625
2012	335	2018	492	2024	215
Total, 2007 to 2024: 3,967 cases					

**Visualizing the 200ft Statutory Protest Buffer
(Empirical TCAD Parcel Boundaries for Case C14-2007-0131)**




Figure 2: 200ft parcel buffer geometries used for TCAD appraisal aggregation and defining the petition buffer. *Generated: 2026-04-13*

3.2 Case Universe and Sample Selection

The Austin Open Data Portal (data.austintexas.gov), accessed via the Socrata Open Data API (SODA), records all administrative zoning activity, including both substantive rezoning proposals and minor administrative variances, such as 5-foot setback adjustments. Because the petition mechanism under study, formal protest petitions under Texas LGC Chapter 211, does not apply to minor variances, and because extraterritorial jurisdiction (ETJ) properties lack the legal protest-petition mechanism, these cases are excluded.

Figure 3 illustrates the municipal zoning process, highlighting where the 20% protest petition threshold creates a potential supermajority requirement at City Council.




Figure 3: Austin municipal zoning process. The figure schematically shows where a qualifying protest petition could make a three-fourths council vote relevant under the study-period framework. A petition representing 20% of the area of lots immediately adjoining the proposed change would require, under the applicable legal framework, a supermajority vote of the City Council for approval, though recusal or absence of council members can alter the exact number of affirmative votes needed.

Table 2: Historical Panel Descriptive Statistics

Statistic	Gross Site Area (Acres)	Requested Height Delta (ft)	Requested FAR Delta	Requested Bldg Cov Delta (%)	Filing Year	Organized Opposition (Binary)
Count	534	3994	3994	3994	7074	7153
Mean	4.92	10.5	0.43	12.3	2008.6	0.06
Std. Dev	30.66	34.9	0.95	20.7	10.11	0.23
Min	0.00	-60	-3.00	-55	1966	0
Median	0.00	0	0.05	5	2009	0
Max	669.83	365	7.60	80	2026	1

Table 3 reports the sample selection from raw administrative records to the analytic sample.

Table 3: Sample Selection from Raw Records to Analytic Sample

Selection Step	Cases Remaining
Raw Austin open data portal zoning records (2007 to 2024)	15,238
Restrict to discretionary cases (rezonings, PUDs, NPAs)	7,153
Exclude cases with unusable or missing initial filing-year information	6783
Final Analytic Sample (Stages C/D)	6783

The analytical samples vary by analytical goal. The primary predictive evaluation, the filing-date protest-petition model, uses the full sample of 6783 discretionary geographic cases from the legally relevant pre-HB 24 petition regime (2007 to 2024). Supporting analyses are constrained to institutionally defensible windows rather than broad pooled periods. The electoral-covariate descriptive panel ($N = 518$) is limited to active single-member districts over a restricted 2022 to 2024 calendar window and is reported as appendix-only non-identifying evidence. The HOME appendix analysis uses the same restricted district panel and is presented as non-identifying descriptive event tracking only. The Invariant Causal Prediction (ICP) analysis utilizes $N = 1,186$ multi-parcel spatial environments to represent cohesive spatial clusters. Finally, the institutional geographic overlay appendix aggregates policy flags by council term ($N = 20$ macro-observations) and marks designs with insufficient identifying variation as not estimable.

Unlike complete-case spatial intersection approaches, the filing-date model allows partially observed covariates and uses the algorithm’s native handling of missing values to avoid dropping observations. This design choice reduces artificial attrition but does not resolve the underlying missing-data problem: missingness can remain systematically patterned across space, project type, and administrative pathway.

Missingness is therefore treated as an evidentiary limitation and reported descriptively rather than framed as solved by estimator choice.

3.3 Primary Outcome Definition

The primary predictive outcome (y_0) is a *measured threshold-crossing petition*: a binary indicator equal to one when the reconstructed petition-share calculation, derived from public petition filings and buffer-area computations, exceeds the 20% statutory threshold defined by Texas Local Government Code Chapter 211 and Austin LDC § 25-2-284 [58]. This is a defined outcome evaluated from public records, not a direct observation of the legally decisive event. The City Clerk’s formal certification of each petition as legally valid is not recorded in the public record; consequently, the outcome captures cases that *appear* to have crossed the threshold based on available records, which may include petitions that were ultimately rejected on administrative grounds or petitions whose true area share differed from the calculated value.

Terminology. Throughout this manuscript, “measured threshold-crossing petition” refers to the observed binary study outcome, “reconstructed petition share” refers to the computed share of signed adjoining-lot area derived from petition filings and parcel spatiality, and “clerk-certified valid petition (unobserved)” refers to the legal administrative determination by the City Clerk that is not captured in the analytic sample.

Cases withdrawn before public notice are excluded from the sample, rather than coded as zero, because no opportunity for formal public opposition existed. Extraterritorial Jurisdiction (ETJ) cases and routine minor variances are additionally excluded because the statutory petition mechanism does not mathematically apply to them.

3.4 Statutory Context: The Legal Petition Mechanism

Formal protest petitions represent a strict administrative mechanism granting incumbent property owners veto authority over localized zoning changes. As detailed by Furth and McKinley, twenty-one U.S. states actively enforce petition statutes that allow a localized numeric minority of nearby residents to force supermajority voting constraints on municipal legislative bodies [22]. While specific standards vary geographically—for instance, Oklahoma requires a 50% sign-on threshold spanning 300 feet, whereas Iowa utilizes a 20% threshold at 200 feet—the institutional architecture remains remarkably consistent: protest petitions formally subordinate city-wide housing coordination to the exclusionary friction of hyper-localized abutters.

Texas historically enforced one of the nation’s most active 20% protest thresholds, culminating in 2025 legislative reforms that explicitly exempted comprehensive mapping efforts and residential upzonings from the supermajority requirement precisely to curb the immense veto power yielded by neighborhood opposition groups [22]. Figure 4 illustrates a representative historical petition filed for a development rezoning in Austin, documenting the precise computational format through which property owners within a tight 200-foot administrative buffer aggregate formal opposition to legally bind the City Council.

3.5 Feature Libraries: Policy Forecasting vs. Explanatory Research

The feature set is separated to prevent predictive models from leveraging sensitive attributes:

Predictive features include physical site characteristics (lot area, unit count changes, PUD/TOD flags), statutory eligibility flags (HB 24 status, owner-vs.-city-initiated), aggregated buffer economics (homestead shares, land-to-improvement ratios), broad demographic indicators (median income, renter share), historical petition rates within the council district, and macroeconomic variables (mortgage rates, vacancy rates) [19]. These elements are selected strictly based on their availability in public

PETITION				
Case Number:		C14-2007-0131	Date:	July 29, 2008
1506 WALLER ST, 807 E 16TH ST & 908 E 15TH ST				
Total Area Within 200' of Subject Tract		293,002.55		
1	02-0906-0502	MONAHAN CASEY J	7423.26	2.53%
2	02-0906-0503	MONAHAN CASEY J	7390.73	2.52%
		CLARK MICHAEL G &		
3	02-0906-0504	RITA S DE BE	7354.85	2.51%
4	02-0906-0506	MARTINSON KATHY L	1645.43	0.56%
		CHILES ROSALIE		
5	02-0906-0507	BENSON	1618.17	0.55%
6	02-0906-0508	KOCHERT KELLEY	1612.79	0.55%
7	02-0906-0509	MONAHAN CASEY J	1419.72	0.48%
8	02-0906-0606	TIMMERMANN TERRELL	3452.33	1.18%
		MERCADO FREDDY G		
9	02-0906-0607	& MARIA R	9482.82	3.24%
10	02-0906-0901	LEGETT GEORGIA F	4459.36	1.52%
11	02-0906-0911	LEGETT GEORGIA F	7856.05	2.68%
		MEDINA JAMES &		
12	02-0906-1001	KRISTINE M GARA	13204.88	4.51%
13	02-0906-1011	RECER DANALYNN	9454.11	3.23%
14	02-0906-1013	BRINSMADE LOUISA C	7634.37	2.61%
15				0.00%
16				0.00%
17				0.00%
18				0.00%
19				0.00%
20				0.00%
21				0.00%
22				0.00%
23				0.00%
24				0.00%
25				0.00%
26				0.00%
27				0.00%
Validated By:		Total Area of Petitioner:		Total %
Stacy Meeks		84,008.88		28.67%

Figure 4: First page summary of the formal protest petition document filed for a Planned Development rezoning in Austin. The municipal record details the physical tracking of petitioner land ownership as a percentage of the total valid area within 200 feet of the subject tract to validate the 20% supermajority threshold.

Table 4: Illustrative statutory parameters governing protest petitions across states detailed by Furth and McKinley [22]. While 21 states utilize similar architectures, the threshold mechanism and buffer stringency dictate the severity of the institutional veto.

State	Validation Threshold	Proximity Buffer	Approval Override Requirement
Texas	20% of nearby land	200 feet	3/4ths Legislative Supermajority
Iowa	20% of nearby land	200 feet	3/4ths Legislative Supermajority
Oklahoma	20% and 50% tiers	300 feet	Legislative Supermajority
Arizona	20% of land & 20% of lots	Municipal Discretion	Legislative Supermajority
Massachusetts	Veto Restricted	Municipal Discretion	Simple Majority (Upzonings)
Michigan (County)	20% of nearby land	Municipal Discretion	Electoral Referendum

administrative and appraisal databases prior to the initial filing date, ensuring the algorithm’s decisions replicate constraints faced by municipal planning staff in real time.

Explanatory research features are restricted to equity analysis. These include tract-level racial and ethnic composition, segregation indices, and probabilistic surname-based racial classification [56]. These features are strictly partitioned from the predictive model to reduce the risk of encoding disparate impacts, though proxy correlation through permitted features remains a concern (see Section 8). By reserving these variables exclusively for post-hoc fairness evaluations, the model avoids generating predictions based on protected class attributes while maintaining the capability to measure intersectional disparities across geographic subgroups.

3.6 Illustrative Panel Structure

The following table illustrates a representative cross-section of the core analytic panel. Each row constitutes a unique discretionary zoning case, enriched with spatial network indicators (historical petition clustering), demographic variables (ACS), and property tax valuations (TCAD EARS) strictly observable before the initial filing date. These specific columns are highlighted because they ultimately emerge as the primary predictive drivers in the meta-attribution analysis.

Case Number	Year	Lag	Petitions (1km)	Median Income	Owner Share	Appraised Value
C14-2015-001	2015	3		\$85,000	62%	\$450,000
C14-2015-002	2015	0		\$42,000	21%	\$210,000
C14-2015-006	2015	7		\$112,000	85%	\$890,000
C14-2016-009	2016	1		\$55,000	34%	\$315,000
C14-2016-012	2016	4		\$95,000	71%	\$620,000

4 Modeling Strategy

4.1 Formal Protest Petition Model

The Formal Protest Petition forecast is the core of the thesis: conditional on a project entering the formal public process, predicting the probability of a measured threshold-crossing petition ($y_0 = 1$):

$$P(y_0 = 1 \mid D = 1, X_{i,t}, Z) \quad (1)$$

At the initial filing date horizon, the predictive feature space ($X_{i,t}$) is restricted to information observable by planners at the time of filing, including site geometry, requested zoning changes, buffer

demographics, and historical petition rates. The primary filing-date specification is intentionally unweighted. Precision-recall curves comparing the CatBoost, Random Forest, ElasticNet Logistic, and V-REx models are reported in Figure 5.

4.1.1 Primary Model Selection

The primary model chosen for all filing-date analyses is **CatBoost** because the prespecified evaluation criteria prioritizes out-of-sample (temporal holdout) PR-AUC, subject to a constraint on probability reliability (lowest ECE among highly ranked models).

4.2 Model Evaluation Strategy

Precision-Recall AUC (PR-AUC) is the primary discrimination metric, given the substantial class imbalance of the outcome. Following standard methodological practices in rare-event spatial forecasting, traditional Receiver Operating Characteristic (AUROC) metrics are omitted entirely, as the large number of true negatives artificially inflates the curve. PR-AUC explicitly isolates the minority class, forcing the model to balance recall against the precision cost of false alarms.

Primary evaluation. The out-of-distribution bootstrap PR-AUC at the filing-date horizon is the primary evaluation object, as it approximates forward-like administrative use. In-distribution 5-fold cross-validation is retained only as an optimistic upper-bound object. Because this is primarily a ranking exercise, emphasis remains on ranking discrimination (PR-AUC, top-decile lift), with calibration and thresholded diagnostics treated as supporting information rather than coequal headline estimands.

The mean predicted probability for true positive cases is 0.480, compared to 0.008 for true negatives. All performance benchmarking relies on strictly temporal constraints (rolling-origin splits). This guarantees the algorithms are evaluated on their ability to predict future, unseen market conditions using only historically available data, directly simulating a real-world historical scenario.

Council composition and policy conditions changed over the sample window, but the temporal analysis does not treat any single election as an exogenous policy break. Performance differences across years are reported descriptively as out-of-sample temporal variation.

4.3 Outcome Scope

This thesis strictly predicts whether a neighborhood protest petition will be filed, rather than predicting the final City Council up-or-down vote on the development. We avoid predicting final council votes because that data is fundamentally incomplete: developers frequently withdraw their zoning applications early if they face heavy neighborhood backlash. By focusing entirely on the initial protest petition filing, the analysis targets a clean, measurable indicator of neighborhood opposition before behind-the-scenes negotiations or case withdrawals permanently hide the outcome.

4.4 Feature Attribution and Semantic Clustering

To evaluate the predictive mechanisms driving protest risk, the analysis calculates feature importance using SHAP (SHapley Additive exPlanations) values. Because administrative datasets contain highly correlated proxies (e.g., tax assessor appraised value, market value, and land value), individual raw features suffer from SHAP attribution dilution where importance weights are split arbitrarily among collinear variables.

To prevent collinearity from obscuring substantive relationships and to ensure the archetypal profiles remain interpretable, the raw pipeline variables are grouped into plain-English “Semantic Clusters” prior

to attribution visualization. Table 5 crosswalks these overarching semantic concepts back to their exact, raw administrative database columns.

Table 5: Semantic Feature Aggregation Mapping. This table crosswalks the plain-English conceptual buckets used in the manuscript’s attribution analysis back to their constituent raw variables from the City of Austin and the completely collinear tax assessor equivalents.

Semantic Cluster	Raw Data Dictionary Variables
Demographic Composition	ACS % White, % Hispanic, % Black, % Asian
Housing Tenure	ACS Owner-Occupied Units, Renter-Occupied Units, Total Units
Nbhd. Income & Rent	ACS Median Household Income, Median Gross Rent, Poverty Count, Median Home Value
Property Valuation	Tax Assessor Appraised Value, Market Value, Land Market Value
Structure Age	Tax Assessor Year Built, COA Year Built
Parcel Scale	Tax Assessor Acres, Lot Size (sqft), Planning Gross Area, Deed Acreage
Land Use Classification	Tax Assessor Land Use, COA Land Use Inventory (LUI)
Zoning Density	Tax Assessor Floor-to-Area Ratio (FAR), Units
Improvement Scale	Tax Assessor Improvement Square Footage
Historical Protest Activity	Binary Prior Protest Flag, 1-Year Spatial Contagion, 3-Year Spatial Contagion
Filing Timeline	Target Filing Year

5 Filing-Date Primary Results

5.1 Predictive Performance Across All Horizons

Because the inherent rarity of protest petitions (a 2.6% base rate) makes absolute probability calibration mechanically misleading, this thesis evaluates performance strictly on relative ranking capacity. Under the primary temporal evaluation, the baseline filing-date model ranks petition risk significantly above random chance, achieving an out-of-distribution PR-AUC of 0.718 (95% CI: [0.48, 0.91]).¹

This ranking utility is best understood through top-decile lift. When isolating the 10% of cases the model flags as highest risk, the precision is 25.4% (a lift of 9.9x over the base rate). However, because there are few true positive cases even in this extreme top decile (15 out of 59 cases), the precision estimate retains a wide confidence interval (16.1% to 37.8%).

These constraints establish the overarching narrative of the predictive results: the algorithmic framework successfully learns to triage neighborhood opposition risk based purely on filing-date administrative geometry and historical spatial friction. However, the absolute scarcity of petitioned cases means that

¹The 95% CI is wide because the temporal holdout split contains only $n_{\text{pos}} = 15$ positives among $n_{\text{test}} = 583$ cases. The interval is computed via bootstrap resampling of held-out test rows.

these scores must remain bounded to relative triage rankings, and cannot be repurposed as deterministic, case-level probabilities for formal case-level results.




Figure 5: Precision-recall curves for filing-date petition models from the pre-registered benchmark roster.

These results provide one transparency diagnostic for feature reliance. The in-distribution cross-validation benchmark is retained as an optimistic upper-bound object in the appendices rather than as a coequal main-text exhibit.

Temporal Predictive Drift (Rolling Origin). To quantify the model’s predictive half-life, Table 26 tracks PR-AUC for each annual window for the benchmark architectures. This is treated as a concept-drift diagnostic rather than a regime-shift design [23]. For each anchor year t , each model is trained on data before t and evaluated forward across all subsequent evaluation years.

While CatBoost frequently offers strong holdout performance, minority architecture wins are concentrated in a subset of years, indicating temporal heterogeneity rather than stable, permanent cross-architecture superiority.

The central temporal finding is a dip-and-recovery pattern: CatBoost’s Pre-2019 anchor decays from 0.659 at immediate horizon (2019, T+0) to 0.329 at the 2022 window (T+3), then rebounds to 0.885 in 2024 (T+5). This is treated as descriptive out-of-sample variation and does not identify a single institutional break mechanism.

Because the absolute PR-AUC baseline (the petition occurrence rate) drops substantially in the later

Table 6: **Temporal predictive drift: PR-AUC decay by algorithm (Optimal Entropy-Based Discretization).**

Model	Anchor Training	2018	2019	2020	2021	2022	2023	2024
CatBoost	Pre-2018	0.960	0.794	0.472	0.324	0.258	0.591	0.622
CatBoost	Pre-2019	—	0.524	0.481	0.405	0.143	0.771	0.395
CatBoost	Pre-2020	—	—	0.503	0.354	0.208	0.636	0.430
CatBoost	Pre-2021	—	—	—	0.467	0.325	0.799	0.415
CatBoost	Pre-2022	—	—	—	—	0.625	0.367	0.326
CatBoost	Pre-2023	—	—	—	—	—	0.633	0.443
CatBoost	Pre-2024	—	—	—	—	—	—	0.586
XGBoost	Pre-2018	0.926	0.553	0.467	0.341	0.208	0.485	0.579
XGBoost	Pre-2019	—	0.397	0.399	0.368	0.208	0.368	0.679
XGBoost	Pre-2020	—	—	0.511	0.276	0.196	0.543	0.458
XGBoost	Pre-2021	—	—	—	0.402	0.267	0.618	0.679
XGBoost	Pre-2022	—	—	—	—	0.250	0.443	0.679
XGBoost	Pre-2023	—	—	—	—	—	0.525	0.350
XGBoost	Pre-2024	—	—	—	—	—	—	0.442
Logistic (L2)	Pre-2018	0.852	0.700	0.413	0.236	0.139	0.277	0.771
Logistic (L2)	Pre-2019	—	0.667	0.358	0.280	0.132	0.244	0.360
Logistic (L2)	Pre-2020	—	—	0.347	0.281	0.133	0.302	0.367
Logistic (L2)	Pre-2021	—	—	—	0.264	0.162	0.258	0.278
Logistic (L2)	Pre-2022	—	—	—	—	0.202	0.348	0.296
Logistic (L2)	Pre-2023	—	—	—	—	—	0.360	0.320
Logistic (L2)	Pre-2024	—	—	—	—	—	—	0.305
Spatial-FE Logistic	Pre-2018	0.642	0.603	0.387	0.296	0.113	0.362	1.000
Spatial-FE Logistic	Pre-2019	—	0.694	0.403	0.341	0.127	0.416	0.374
Spatial-FE Logistic	Pre-2020	—	—	0.440	0.334	0.110	0.420	0.379
Spatial-FE Logistic	Pre-2021	—	—	—	0.285	0.134	0.304	0.305
Spatial-FE Logistic	Pre-2022	—	—	—	—	0.138	0.342	0.287
Spatial-FE Logistic	Pre-2023	—	—	—	—	—	0.471	0.298
Spatial-FE Logistic	Pre-2024	—	—	—	—	—	—	0.306
TabNet	Pre-2018	0.990	0.399	0.802	0.226	0.267	0.295	0.380
TabNet	Pre-2019	—	0.364	0.506	0.373	0.174	0.273	0.351
TabNet	Pre-2020	—	—	0.309	0.235	0.167	0.173	0.170
TabNet	Pre-2021	—	—	—	0.503	0.244	0.454	0.775
TabNet	Pre-2022	—	—	—	—	0.216	0.349	0.461
TabNet	Pre-2023	—	—	—	—	—	0.254	0.390
TabNet	Pre-2024	—	—	—	—	—	—	0.667
Anchor Regression (Causal)	Pre-2018	0.994	0.452	0.500	0.432	0.133	0.343	0.525
Anchor Regression (Causal)	Pre-2019	—	0.539	0.500	0.314	0.171	0.237	0.406
Anchor Regression (Causal)	Pre-2020	—	—	0.273	0.331	0.208	0.230	0.661
Anchor Regression (Causal)	Pre-2021	—	—	—	0.331	0.267	0.303	0.575
Anchor Regression (Causal)	Pre-2022	—	—	—	—	0.155	0.427	0.722
Anchor Regression (Causal)	Pre-2023	—	—	—	—	—	0.356	0.469
Anchor Regression (Causal)	Pre-2024	—	—	—	—	—	—	0.557

years of the evaluation period, absolute PR-AUC mechanically declines even if relative ranking separation remains perfectly stable. To isolate this baseline-shift effect, Table 27 reproduces the temporal drift analysis measuring relative PR-AUC lift (dividing the absolute PR-AUC across the entire threshold curve by the contemporaneous evaluation-year base rate) to measure dynamic model intelligence relative to random guessing.

Normalizing against the collapsing base rate reveals a profoundly counter-intuitive *signal dilution* effect. First, models trained exclusively on the high-conflict Pre-2018 era generate a blistering $34\times$ lift multiplier when predicting the "dead" legislative environment of 2024. Because the spatial recipe for formal opposition was hardened into the high-volume historical data, freezing the legacy model mathematically preserves the spatial sorting intelligence, allowing it to isolate the handful of surviving post-ordinance protests with extreme precision relative to random chance. However, iteratively retraining the model on modern data (e.g., Pre-2023 or Pre-2024 anchors) paradoxically *degrades* the 2024 out-of-sample lift by nearly half (dropping to $\sim 19\times$). Exposing the algorithm to the post-ordinance era, where the true-positive density functionally evaporates, starves the model of positive labels and induces *concept drift*. The algorithm stops severely penalizing aggressive spatial configurations because the silenced contemporary institutional environment has stopped explicitly penalizing them.

While aggregate architectural performance appears statistically indistinguishable, grouping architectures by functional capacity reveals strict boundaries around domain shift performance (Table 28). First, linear architectures (Logistic Regression, L2) exhibit extreme *horizon attrition*, demonstrating competitive dominance almost exclusively at immediate-year extrapolations (T+0, T+1) before decaying abruptly at distant forecasting horizons. Second, non-linear tree ensembles exhibit superior resilience to *origin volatility*. When training anchors are forced to bridge the institutional break characterizing the post-ordinance era (Pre-2020 through Pre-2024 anchors), boosted trees dominate 15 out of 16 out-of-sample evaluation cells, confirming that non-linear partition mapping is algorithmically required to parse highly volatile institutional environments.

5.2 Predictors of Petition Filing

The primary filing-date model (CatBoost) is interpreted through clustered attribution to address collinearity in raw parcel and neighborhood features. Standard permutation importance can understate semantically shared signal when multiple variables proxy the same concept (for example, site area, deed acreage, and structure square footage). Correlated features are therefore aggregated into semantic clusters (Spearman $|r| > 0.70$) and summed relative importance is reported by cluster (Figure 8; full feature-to-cluster mapping in Appendix ??).² For the primary model, reliance concentrates on parcel scale, property valuation, and neighborhood composition.

Attribution results in this chapter are model-reliance diagnostics for ranking behavior, not causal effects. Full case-level SHAP beeswarm decompositions are reported in Appendix B.3 [38, 37].

5.3 Attribution Stability and Meta-Attribution


Meta-attribution is the main transparency audit in this section. Figure 9 presents a Ward-linkage clustermap comparing architectures and origin years, showing a consistent functional split: tree-based models and deep models occupy different attribution regimes. The main-text claim is architecture dependence, not a discrete regime-shift claim. Architecture choice changes which feature families drive

²A semantic ablation test confirmed that pre-clustering features into conceptual averages prior to training leads to substantial information loss, particularly regarding parcel-scale drivers. Using post-hoc aggregation of raw TreeSHAP values preserves model-level fidelity while improving interpretability under collinearity [38, 37].

Table 7: **Temporal predictive drift: PR-AUC lift by algorithm (Optimal Entropy-Based Discretization).**

Model	Anchor Training	2018	2019	2020	2021	2022	2023	2024
CatBoost	Pre-2018	6.307	11.904	8.346	6.944	27.818	29.693	33.407
CatBoost	Pre-2019	—	7.861	8.512	8.671	15.429	38.734	21.212
CatBoost	Pre-2020	—	—	8.910	7.579	22.500	31.965	23.132
CatBoost	Pre-2021	—	—	—	10.000	35.100	40.130	22.332
CatBoost	Pre-2022	—	—	—	—	67.500	18.460	17.538
CatBoost	Pre-2023	—	—	—	—	—	31.825	23.825
CatBoost	Pre-2024	—	—	—	—	—	—	31.503
XGBoost	Pre-2018	6.084	8.301	8.267	7.302	22.500	24.384	31.130
XGBoost	Pre-2019	—	5.961	7.069	7.871	22.500	18.497	36.505
XGBoost	Pre-2020	—	—	9.053	5.909	21.214	27.279	24.635
XGBoost	Pre-2021	—	—	—	8.613	28.800	31.047	36.505
XGBoost	Pre-2022	—	—	—	—	27.000	22.270	36.505
XGBoost	Pre-2023	—	—	—	—	—	26.381	18.839
XGBoost	Pre-2024	—	—	—	—	—	—	23.740
Logistic (L2)	Pre-2018	5.595	10.494	7.317	5.040	15.058	13.895	41.432
Logistic (L2)	Pre-2019	—	10.001	6.345	5.997	14.308	12.240	19.330
Logistic (L2)	Pre-2020	—	—	6.149	6.003	14.400	15.172	19.708
Logistic (L2)	Pre-2021	—	—	—	5.648	17.532	12.945	14.964
Logistic (L2)	Pre-2022	—	—	—	—	21.808	17.475	15.901
Logistic (L2)	Pre-2023	—	—	—	—	—	18.082	17.204
Logistic (L2)	Pre-2024	—	—	—	—	—	—	16.381
Spatial-FE Logistic	Pre-2018	4.218	9.048	6.859	6.338	12.214	18.184	53.750
Spatial-FE Logistic	Pre-2019	—	10.412	7.139	7.306	13.708	20.905	20.116
Spatial-FE Logistic	Pre-2020	—	—	7.794	7.138	11.868	21.117	20.372
Spatial-FE Logistic	Pre-2021	—	—	—	6.092	14.464	15.282	16.381
Spatial-FE Logistic	Pre-2022	—	—	—	—	14.914	17.199	15.444
Spatial-FE Logistic	Pre-2023	—	—	—	—	—	23.689	16.004
Spatial-FE Logistic	Pre-2024	—	—	—	—	—	—	16.452
TabNet	Pre-2018	6.498	5.979	14.191	4.837	28.800	14.814	20.409
TabNet	Pre-2019	—	5.458	8.952	7.985	18.750	13.706	18.850
TabNet	Pre-2020	—	—	5.471	5.037	18.000	8.687	9.113
TabNet	Pre-2021	—	—	—	10.758	26.308	22.792	41.656
TabNet	Pre-2022	—	—	—	—	23.318	17.513	24.785
TabNet	Pre-2023	—	—	—	—	—	12.753	20.988
TabNet	Pre-2024	—	—	—	—	—	—	35.833
Anchor Regression (Causal)	Pre-2018	6.526	6.774	8.850	9.241	14.400	17.229	28.219
Anchor Regression (Causal)	Pre-2019	—	8.087	8.850	6.718	18.514	11.934	21.799
Anchor Regression (Causal)	Pre-2020	—	—	4.832	7.088	22.500	11.556	35.513
Anchor Regression (Causal)	Pre-2021	—	—	—	7.092	28.800	15.205	30.906
Anchor Regression (Causal)	Pre-2022	—	—	—	—	16.714	21.461	38.819
Anchor Regression (Causal)	Pre-2023	—	—	—	—	—	17.877	25.233
Anchor Regression (Causal)	Pre-2024	—	—	—	—	—	—	29.946

Temporal Predictive Drift by Model Architecture (PR-AUC)



Temporal Predictive Drift by Model Architecture (Lift)





Figure 6: **Temporal drift evidence across terminal observation years.** Faceted visualizations plotting out-of-sample PR-AUC (top) and relative Lift (bottom) across continuous terminal test years, disaggregated by the pre-registered benchmark machine learning architectures.

Temporal Predictive Drift by Model Architecture (PR-AUC)



Temporal Predictive Drift by Model Architecture (Lift)




Figure 7: **Temporal drift evidence relative to Out-Of-Distribution temporal offset.** Faceted visualizations plotting out-of-sample PR-AUC (top) and relative Lift (bottom) tracked against elapsed time offsets (+Years) from the original anchor training boundary, summarizing architecture-dependent and origin-dependent forecasting degradation. 23

Table 8: Max-of-Family dominance (Optimal Entropy-Based Discretization).

Anchor Training	2018	2019	2020	2021	2022	2023	2024
Panel A: Maximum Absolute PR-AUC							
Pre-2018	0.994 (Causal)	0.794 (Trees)	0.802 (Deep)	0.432 (Causal)	0.267 (Deep)	0.591 (Trees)	1.000 (Linear)
Pre-2019	—	0.694 (Linear)	0.506 (Deep)	0.405 (Trees)	0.208 (Trees)	0.771 (Trees)	0.679 (Trees)
Pre-2020	—	—	0.511 (Trees)	0.354 (Trees)	0.208 (Causal)	0.636 (Trees)	0.661 (Causal)
Pre-2021	—	—	—	0.503 (Deep)	0.325 (Trees)	0.799 (Trees)	0.775 (Deep)
Pre-2022	—	—	—	—	0.625 (Trees)	0.443 (Trees)	0.722 (Causal)
Pre-2023	—	—	—	—	—	0.633 (Trees)	0.469 (Causal)
Pre-2024	—	—	—	—	—	—	0.667 (Deep)
Panel B: Maximum Relative PR-AUC Lift							
Pre-2018	6.53 (Causal)	11.90 (Trees)	14.19 (Deep)	9.24 (Causal)	28.80 (Deep)	29.69 (Trees)	53.75 (Linear)
Pre-2019	—	10.41 (Linear)	8.95 (Deep)	8.67 (Trees)	22.50 (Trees)	38.73 (Trees)	36.51 (Trees)
Pre-2020	—	—	9.05 (Trees)	7.58 (Trees)	22.50 (Causal)	31.96 (Trees)	35.51 (Causal)
Pre-2021	—	—	—	10.76 (Deep)	35.10 (Trees)	40.13 (Trees)	41.66 (Deep)
Pre-2022	—	—	—	—	67.50 (Trees)	22.27 (Trees)	38.82 (Causal)
Pre-2023	—	—	—	—	—	31.82 (Trees)	25.23 (Causal)
Pre-2024	—	—	—	—	—	—	35.83 (Deep)

Comparative Primary Reliance: Top Feature Clusters Across Architectures




Figure 8: Top semantic predictor groups aggregated by average reliance within the four formalized model families (Tree, Deep, Regularized Linear, Distributionally Robust). Distinct reliance regimes hold uniquely between computational architectures: Linear and Deep proxies track spatial scale (Land Acres), while Tree methods locate heavy variance locally within demographics and zoning density.

ranking behavior, even when predictive tasks and data are held constant. The heatmap itself suggests a recurring socio-demographic reliance band rather than a complete absence of overlap: demographics is repeatedly the largest or near-largest cluster, neighborhood income and housing tenure recur as the next tier in most model-year cells, and parcel or property valuation features enter as secondary support. What varies by architecture is the proportional mix, especially the relative emphasis on social-context versus physical-value feature clusters, not whether any overlap exists at all.

To explicitly quantify the divergence mapped in Figure 9, Table 9 aggregates the raw clustered reliance weights into archetypal averages by architectural family. Tree ensembles allocate over 80% of their predictive logic exclusively to socio-demographic contexts (Demographics, Housing Tenure, Income), practically ignoring physical parcel spatiality. Deep architectures heavily re-weight this attention, slashing demographic reliance and focusing intensely on physical and geographic attributes (Parcel Scale, Property Valuation, and Structure Age).

Crucially, to prevent the model from memorizing parcel identity (e.g., exploiting high-cardinality exact physical floats such as an isolated \$1,219,301 market valuation as unique identification hashes), the unified 4-architecture benchmark discretizes continuous physical features using entropy-based binning ($\text{min_samples_leaf} = 100$). Rather than relying on arbitrary quantiles (e.g., deciles) or theoretically fragile uniform-width rules (like the Freedman-Diaconis heuristic), this approach ensures that no resulting bin can ever contain fewer than 100 properties. Consequently, the continuous uniqueness of extreme outliers is prevented from being exploited, ensuring that all 4 baseline architectures are evaluated strictly on their ability to learn generalized, reproducible spatial concepts rather than discrete memorization boundaries.

Table 9: **Archetypal Family Attribution.** Average absolute model reliance allocated to each semantic feature cluster. **Causal** array maps the Non-Linear Anchor Regression. High-cardinality identity floats mathematically blinded via Shannon-Entropy decision tree mapping ($\text{min_samples_leaf} = 100$).


Semantic Target Cluster	Tree Ensembles	Deep Architectures	Linear Architectures	Causal Algorithms
Parcel Scale	40.5%	15.7%	19.4%	21.7%
Demographics	20.8%	13.5%	22.8%	18.3%
Structure Age	10.9%	12.4%	2.9%	8.7%
Zoning Density	8.6%	11.3%	8.2%	14.3%
Neighborhood Income	6.9%	4.0%	10.7%	9.1%
Property Valuation	6.2%	29.4%	4.3%	2.8%
Housing Tenure	3.7%	10.0%	14.6%	8.9%
Improvement Scale	2.3%	1.7%	9.9%	8.3%
Neighborhood Valuation	0.2%	2.1%	7.2%	7.9%

We extend this analysis by weighting the standard attribution space relative to actual OOS robustness. Table 10 maps the performance-weighted archetypal boundaries, wherein highly performant, robust configurations dominate the family mean, answering exactly what mathematically valid models depend on to degrade gracefully under longitudinal drift.

Temporal attribution persistence remains fundamentally coupled to architectural geometry. Figure 12 generalizes this window to the four meta-families, computing sustained cross-year adjacent-window Spearman rank correlation stability under shifting observation spans.

Operationally, the attribution sidecar supports a narrower, more qualified conclusion: there is a recurring set of high-weight clusters, especially demographics, neighborhood income, and housing tenure, but their relative shares shift enough between architectures and years that they should be treated as a

Meta-Attribution Structural Clustering



Performance-Weighted Meta-Attribution Clustering




Table 10: **Performance-Weighted Archetypal Family Attribution.** Model attribution vectors scaled by their corresponding out-of-distribution longitudinal PR-AUC performance. Architectures that successfully generalized under domain drift dominate their archetypal class weights.

Semantic Target Cluster	Tree Ensembles	Deep Architectures	Linear Architectures	Causal Algorithms
Parcel Scale	42.3%	14.7%	19.7%	20.8%
Demographics	20.5%	13.3%	22.8%	18.9%
Structure Age	9.7%	10.2%	2.8%	8.6%
Zoning Density	8.2%	11.5%	8.4%	15.2%
Neighborhood Income	6.6%	2.2%	10.3%	10.1%
Property Valuation	6.4%	33.4%	4.7%	2.8%
Housing Tenure	4.0%	9.1%	14.2%	8.7%
Improvement Scale	2.3%	2.6%	10.2%	8.0%
Neighborhood Valuation	0.1%	3.0%	7.0%	6.9%




Figure 11: Attribution Stability: Continuous Divergence Tracing. Year-by-year longitudinal tracking from 2018 to 2024 comparing internal model reliance on Demographics versus Parcel Scale between the four unified benchmark architectures. Deep and Linear proxies cleanly fixate continuous spatial topology, while Tree models permanently lock into socio-demographic splitting variance in otherwise identical spatial matrices.




Figure 12: Comprehensive architectural stability map reporting the adjacent-window Spearman rank correlation (ρ) of internal semantic feature ordering natively calculated for identically trained Tree, Deep, Regularized Linear, and Distributionally Robust baseline architectures throughout the 2018–2024 rolling windows.

descriptive pattern rather than a fixed portable explanatory profile. The practical use of the attribution sidecar is therefore audit and monitoring, not a claim that any named feature cluster travels unchanged across model classes. Full SHAP beeswarm exhibits are reported in Appendix B.3.

6 Institutional Context

This chapter presents bounded institutional evidence that contextualizes the filing-date petition results. None of the analyses here constitute formal causal identification of the petition mechanism. The reduced-form threshold estimate characterizes the association between measured threshold-crossing and council processing time; electoral-covariate diagnostics are descriptive appendix checks only. Additional overlay, HOME-tracking, and ICP diagnostics are reported in the appendices as non-headline transparency checks.

6.1 Threshold-Based Reduced-Form Discontinuity on Reconstructed Petition Share

To estimate the threshold-associated shift in processing timelines, the analysis applies a threshold-based reduced-form discontinuity on reconstructed petition share at the 20% statutory margin under Texas LGC Chapter 211 [58]. The running variable is the reconstructed percentage of signed area within the 200-foot buffer [36].

A triangular-kernel weighted least squares estimator indicates a positive discontinuity in council processing time at the measured threshold. Full numeric estimates are reported in registry-linked tables; this analysis is treated as supporting institutional context rather than a coequal headline estimand.

This estimate should be interpreted as a reduced-form effect at the threshold. Because the data do not record whether the City Clerk formally certified each petition as legally compliant, the analysis relies on an unverified calculated area share. Rather than assuming the exact threshold acts as a deterministic sharp trigger, the specification estimates the shift in average processing time associated with crossing the

statutory 20% margin. Consequently, this is a threshold-based reduced-form model where institutional compliance remains unobserved.

The identification schematic is reported in Appendix C.1 (Figure 17); the main text retains only the estimated discontinuity exhibit.

Design limitations include potential manipulation of the running variable, since petition organizers may strategically target just above 20%, and limited local support near the cutoff in some calendar slices. Robustness checks, a McCrary (2008) density test for running-variable manipulation [42] ($p > 0.05$), symmetric donut exclusions, and covariate continuity checks, support the validity of the design at the threshold.




Figure 13: Threshold-based reduced-form discontinuity at the 20% reconstructed petition-share threshold. The figure reports the estimated shift in council processing time. Bin widths were selected using the standard IMSE-optimal (Integrated Mean Squared Error) procedure [9].

7 Qualitative Planning Context

The quantitative results and interview interpretations are reported as distinct forms of evidence. Quantitative models establish empirical claims about predictive discrimination, calibration, and temporal performance variation; interviews provide institutional context for how those patterns are experienced in practice. Semi-structured interviews with urban planning professionals and city staff (conducted under Columbia University IRB Protocol ACYY0820) surface dynamics that administrative data cannot capture. The interview sample is purposive, small- N , and not statistically representative of all Austin stakeholders. Interview material is therefore used for interpretive contextualization only and is not used to estimate effect sizes, test hypotheses, or justify formal case-level decisions.

The quantitative findings that the interviews are intended to contextualize are straightforward:

- Temporal performance varies over holdout years, with a 2021 to 2022 dip and later recovery in filing-date discrimination.
- Petition-risk predictions retain ranking utility, but wide PR-AUC uncertainty and sparse positives still constrain formal use of case-level probability estimates despite a numerically small ECE on the evaluated score column ($ECE = 0.009$).
- Exploratory electoral-covariate diagnostics are non-identifying, appendix-only checks and are not used to support causal or regime-shift claims.

Interview-based interpretation of institutional context follows:

- **Perceived decline in petition leverage over time:** Interviewees described reduced practical leverage of petitions in the later study period and emphasized broader institutional evolution rather than a single-date break. Annick Beaudet (former zoning planner) noted that petitions appeared to lose traction during the Land Development Code revision period: “While people still scream and yell, they don’t get as much traction... Necessity is the mother of invention” [3, 13]. These perceptions are qualitative context and are not used as identifying evidence in the quantitative design.
- **The Function of Opposition:** Highlighting the nuanced motivations driving public testimony, AURA board member Felicity Maxwell observed: “We’re not talking about suppressing opposition, but understanding it better to have more productive conversations” [40].
- **Asymmetric Transparency and Silent Supporters:** Jonathan Tomko (Principal Planner) emphasized that the municipal public comment process lacks any mechanism for *anonymous* public testimony against a zoning case [60]. Because neighborhood social dynamics often penalize residents who deviate from their neighborhood association’s preservationist stance, Tomko noted that the formal opposition data is not strictly representative of the entire community. This aligns with Beaudet’s observation that “haters are overrepresented and silent supporters are underrepresented” [3], which likely makes the observed opposition data an incomplete measure of community sentiment.
- **Two distinct motivations and geography:** Marla Torrado (Displacement Prevention Division) stressed that opposition mechanisms may be driven by differing underlying concerns [16]. While the “vocal opposition” on the affluent West side leverage environmental overlays, East Austin neighborhoods frequently mobilize opposition mechanisms in response to rapid gentrification pressures. This context reinforces why rigid probability thresholds are ethically hazardous even when scalar calibration metrics look acceptable on paper: using rigid probability cutoffs could miss communities facing real displacement risk, failing “to identify where people that are in actual need actually are” because the algorithm cannot separate affluent NIMBYism from neighborhood preservation efforts.

These interview interpretations contextualize, rather than establish, the quantitative model constraint: independently of interview evidence, the model is treated as a relative ranking heuristic because discrimination uncertainty and methodological ethics remain limiting even when the evaluated ECE is numerically small.

8 Limitations

This study’s methodological and empirical boundaries introduce several methodological limitations that must contextualize any future application of these predictive tools.

Unobserved Selection and Applicant Anticipation. The analytic dataset fundamentally suffers from unobserved selection bias because it only contains development proposals that formally entered the municipal zoning process. The model cannot observe "chilling effects" where applicants proactively withdrew, heavily modified, or entirely avoided proposing projects in specific neighborhoods because they anticipated severe political hostility. Attempted development-occurrence inverse-probability weighting failed basic overlap and balance diagnostics during preliminary tests and was therefore discarded. Consequently, the baseline opposition rates and model evaluations remain strictly conditioned on applicants actually deciding to file, meaning the algorithm is predicting opposition on an already-censored subset of land use activities.

Measurement Error in Temporal Reconstruction. The thesis relies on reconstructing historical "as-of" dates to prevent future-leakage, forcing the model to rely solely on information that would have been public prior to a petition filing. However, because municipal IT systems regularly overwrite historical tables, administrative milestone dates had to be imputed based on statutory processing minimums. These approximations inherently introduce measurement error into the temporal bounds of the feature matrix, meaning the simulated timing of when neighborhood intelligence became available to organizers is an approximation rather than a deterministic log.

Macro-Level Aggregation and Proxy Risk. The predictive matrix explicitly excludes individual-level demographic profiling (such as Bayesian Improved Surname Geocoding) to limit the weaponization of protected class attributes. Instead, the model depends on macro-level geographic aggregates at the census tract or block group level. Because the predictive engine uses these coarse neighborhood covariates, it cannot resolve fine-grained, behavioral heterogeneity within the immediate 200-foot buffer of a specific parcel. Consequently, the reliance on broad socio-demographic clusters inextricably weaves proxy risk into the model’s architecture, remaining a permanent methodological limitation rather than a solved design feature [2, 5, 46].

Ethical Ambiguity in Algorithmic Error. While the thesis presents formal spatial error distributions, neither geographic symmetry nor false-negative parity constitutes a robust fairness guarantee when algorithms are deployed into politically charged housing markets. The most substantive ethical limitation is that the algorithm’s mathematical architecture is completely blind to motive: it cannot distinguish between affluent, exclusionary NIMBY mobilization intended to block multi-family housing and legitimate, community-driven anti-displacement resistance seeking to protect vulnerable tenants from predatory development (as detailed in Section 7).

External Validity and Policy Windows. Crucially, Austin’s unique institutional structure—characterized by a council-manager government [12], highly localized Texas protest petition statutes, and an archaic continuous land development code—severely limits the direct generalization of these specific predictive parameters to other municipalities. Furthermore, the sweeping Texas HB 24 legislation, which recently gutted the formal petition mechanism, falls entirely outside the thesis observation period. While the overarching computational approach (time-stamped feature engineering, strict sequential evaluation, and clustered attribution tracing) is theoretically portable to other cities, the specific model weights, architectures, and predictive scores presented here are strictly Austin-specific and bounded to a now-closed policy era.

Causal Pathway Limitations and Observational Assumptions. In establishing a causal mediation framework for the "Density Bonus Paradox," this study must adopt several structural assumptions due to the constraints of observational data. First, the *Anchor Pricing Assumption* posits that a developer’s initially requested height or density is genuine; if developers strategically "over-ask"

anticipating negotiations, observed reductions may reflect political theater rather than a true penalizing effect. Second, the *Capital Stack Heterogeneity* assumption acknowledges that not all developers are equally exposed to macroeconomic attrition (e.g., floating-rate debt); while the model proxies institutional, patient capital via PUD or high-density requests, precise internal pro-forma data remains unobserved. Third, the *Qualitative Void Limitation* requires treating discrete hearing counts as a proxy for political friction, mathematically equating a routine 5-minute postponement with a highly contested 4-hour hearing. However, while our current model relies on discrete hearing counts, the pipeline is architected to support future Natural Language Processing (NLP). A pilot NLP analysis of 14,999 commission transcripts confirms this variance, with agenda documentation lengths ranging from 0 to over 166,000 tokens per case, containing hundreds of localized keyword hits for terms like "traffic," "density," and "oppose." By calculating the token count or page-length of the agenda backup materials, future iterations can perfectly proxy the qualitative intensity of the political friction. Finally, *Strategic Withdrawals* represent severe, pre-emptive opposition success that fundamentally censors the final disposition data, requiring the model to rely heavily on earlier protest-petition and hearing delays as the primary visible indicators of political friction.

9 Conclusion

Filing-date administrative features can rank petition risk in Austin’s pre-HB 24 zoning regime, but wide discrimination uncertainty and sparse positives still limit institutionally portable use of case-level probabilities beyond ranking workflows.

This thesis asks whether measured threshold-crossing petitions against housing development proposals can be predicted before formal public hearings. The answer is specific: filing-date administrative features carry real predictive signal for petition risk—the model achieves a PR-AUC of 0.718 on a temporal holdout—but that signal suffered significant temporal degradation during the 2022 performance dip. This suggests that predictive reliability is sensitive to ambient macro-shocks (such as pandemic-era administrative shifts) rather than static parcel spatiality. Furthermore, performance varies erratically across petition thresholds (0%–25%), providing a substantive empirical finding on **neighborhood agency**: neighbor mobilization is a deliberate social choice that subverts mechanistic administrative prediction, confirming that “organized” opposition remains a contingent social process rather than a rule-based response to statutory thresholds. These results demonstrate that while administrative data carries signal, the predictive patterns are mathematical artifacts of a specific policy window and should not be misinterpreted as permanent fundamental laws of neighborhood behavior. Finally, the model’s dependence on socio-demographic clusters—especially neighborhood income, demographics, and housing tenure—indicates that predictive models of petition risk invariably embed complex class and race dynamics. However, because the relative feature importance of these clusters shifts unpredictably between different algorithms and over time, the attribution profiles are not stable enough to extract durable urban planning theories.

Future work should extend the observation window to capture HB 24 effects [57], explore richer text features from pre-hearing public comments, and isolate mechanisms behind temporary covariate shifts over time.

References

- [1] Susan Athey and Guido W. Imbens. Machine learning methods that economists should know about. *Annual Review of Economics*, 11:685–725, 2019. doi:10.1146/annurev-economics-080217-053433.
- [2] Solon Barocas and Andrew D. Selbst. Big data’s disparate impact. *California Law Review*, 104(3):671–732, 2016. doi:10.15779/Z38BG31.
- [3] Annick Beaudet. Interview with annick beaudet, APA-certified urban planner, 2025. Semi-structured interview conducted by Daniel Hardesty Lewis, transcribed via Read.ai.
- [4] Vicki Been, Ingrid Gould Ellen, and Katherine O’Regan. Supply skepticism: Housing supply and affordability. *Housing Policy Debate*, 29(1):25–40, 2019. doi:10.1080/10511482.2018.1476899.
- [5] Ruha Benjamin. *Race After Technology: Abolitionist Tools for the New Jim Code*. Polity Press, 2019.
- [6] Glenn W Brier. Verification of forecasts expressed in terms of probability. *Monthly weather review*, 78(1):1–3, 1950.
- [7] Nadia R. Brouwer and Jessica Trounstein. NIMBYs, YIMBYs, and the politics of land use in American cities. *Annual Review of Political Science*, 27:165–184, 2024. doi:10.1146/annurev-polisci-041322-044444.
- [8] Brantly Callaway and Pedro H.C. Sant’Anna. Difference-in-differences with multiple time periods. *Journal of Econometrics*, 225(2):200–230, 2021. doi:10.1016/j.jeconom.2020.12.001.
- [9] Sebastian Calonico, Matias D. Cattaneo, and Rocío Titiunik. Robust nonparametric confidence intervals for regression-discontinuity designs. *Econometrica*, 82(6):2295–2326, 2014. doi:10.3982/ECTA11757.
- [10] Stephanie Riegg Cellini, Fernando Ferreira, and Jesse Rothstein. The value of school facility investments: Evidence from a dynamic regression discontinuity design. *The Quarterly Journal of Economics*, 125(1):215–261, 2010. doi:10.1162/qjec.2010.125.1.215.
- [11] Carlos Cinelli and Chad Hazlett. Making sense of sensitivity: Extending omitted variable bias. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 82(1):39–67, 2020. doi:10.1111/rssb.12347.
- [12] City of Austin. City of austin charter. Technical report, City of Austin, Austin, TX, 2010. URL: <https://www.austintexas.gov/department/city-charter>.
- [13] City of Austin. CodeNEXT draft 3. Technical report, City of Austin, 2017. URL: <https://www.austintexas.gov/codenext>.
- [14] City of Austin. Resolution no. 20230720-126: Initiating the HOME (Home Options for Mobility and Equity) initiative. Austin City Council, July 2023. Sponsored by Council Member Leslie Pool. Adopted July 20, 2023. URL: <https://www.austintexas.gov/department/home-initiative>.
- [15] City of Austin. Zoning cases and land use inventory datasets. City of Austin Open Data Portal, 2024. URL: <https://data.austintexas.gov>.

- [16] City of Austin Displacement Prevention Division. Interview with marla t., city of Austin displacement prevention division, 2025. Semi-structured interview conducted by Daniel Hardesty Lewis, transcribed via Read.ai.
- [17] Katherine Levine Einstein, David M. Glick, and Maxwell Palmer. *Neighborhood Defenders: Participatory Politics and America's Housing Crisis*. Cambridge University Press, 2019. doi:10.1017/9781108769495.
- [18] Virginia Eubanks. *Automating Inequality: How High-Tech Tools Profile, Police, and Punish the Poor*. St. Martin's Press, 2018.
- [19] Federal Reserve Bank of St. Louis. Federal funds effective rate (FEDFUNDS) and 30-year fixed rate mortgage average (MORTGAGE30US). FRED Economic Data, 2024. URL: <https://fred.stlouisfed.org>.
- [20] William A. Fischel. *The Homevoter Hypothesis: How Home Values Influence Local Government Taxation, School Finance, and Land-Use Policies*. Harvard University Press, 2001. doi:10.4159/9780674036901.
- [21] Salim Furth and C. Whit Ewen. Mostly invisible: The cost of valid petitions in Texas. Policy brief, Mercatus Center at George Mason University, February 2023. URL: <https://www.mercatus.org/research/policy-briefs/mostly-invisible-cost-valid-petitions-texas>.
- [22] Salim Furth and Kelcie McKinley. Rezoning protest petitions are ripe for reform. Policy brief, Mercatus Center at George Mason University, August 2025. URL: <https://www.mercatus.org/housing/policy-briefs/rezoning-protest-petitions-are-ripe-reform>.
- [23] João Gama, Indrė Žliobaitė, Albert Bifet, Mykola Pechenizkiy, and Abdelhamid Bouchachia. A survey on concept drift adaptation. *ACM Computing Surveys*, 46(4):1–37, 2014. doi:10.1145/2523813.
- [24] Edward L. Glaeser and Joseph Gyourko. The impact of building restrictions on housing affordability. *Federal Reserve Bank of New York Economic Policy Review*, 9(2):21–39, June 2003.
- [25] Edward L. Glaeser, Joseph Gyourko, and Raven Saks. Why is Manhattan so expensive? Regulation and the rise in housing prices. *Journal of Law and Economics*, 48(2):331–369, 2005. doi:10.1086/429979.
- [26] Andrew Goodman-Bacon. Difference-in-differences with variation in treatment timing. *Journal of Econometrics*, 225(2):254–277, 2021. doi:10.1016/j.jeconom.2021.03.014.
- [27] Joseph Gyourko, Jonathan Hartley, and Jacob Krimmel. The local residential land use regulatory environment across US housing markets: Evidence from a new Wharton index. *Journal of Urban Economics*, 124:103337, 2021. doi:10.1016/j.jue.2021.103337.
- [28] Joseph Gyourko and Raven Molloy. Regulation and housing supply. In *Handbook of Regional and Urban Economics*, volume 5, pages 1289–1337. Elsevier, 2015. doi:10.1016/B978-0-444-59531-7.00019-3.
- [29] Joseph Gyourko, Albert Saiz, and Anita Summers. A new measure of the local regulatory environment for housing markets: The Wharton residential land use regulatory index. *Urban Studies*, 45(3):693–729, 2008. doi:10.1177/0042098007087341.

- [30] Emily Hamilton and Salim Furth. Housing reform in the states: A menu of options. Policy brief, Mercatus Center at George Mason University, 2023. URL: <https://www.mercatus.org/research/policy-briefs/housing-reform-states>.
- [31] Emily Hamilton and Salim Furth. Texas protest petition primer: How HB 24 reforms the super-majority requirement. Policy brief, Mercatus Center at George Mason University, 2024. URL: <https://www.mercatus.org/research/policy-briefs/texas-protest-petition-primer>.
- [32] Michael Hankinson. When do renters behave like homeowners? high rent, price anxiety, and nimbyism. *American Political Science Review*, 112(3):473–493, 2018. doi:10.1017/S0003055418000035.
- [33] Roderick M. Hills, Jr. and David Schleicher. Balancing the “zoning budget”. *Case Western Reserve Law Review*, 62(1):81–133, 2011.
- [34] Jon Kleinberg, Jens Ludwig, Sendhil Mullainathan, and Ziad Obermeyer. Prediction policy problems. *American Economic Review*, 105(5):491–495, 2015. doi:10.1257/aer.p20151023.
- [35] Koch and Fowler. A city plan for austin, texas. Technical report, City of Austin, 1928.
- [36] David S. Lee and Thomas Lemieux. Regression discontinuity designs in economics. *Journal of Economic Literature*, 48(2):281–355, 2010. doi:10.1257/jel.48.2.281.
- [37] Scott M. Lundberg, Gabriel Erion, Hugh Chen, Alex DeGrave, Jordan M. Prutkin, Bala Nair, Ronit Katz, Jonathan Himmelfarb, Nisha Bansal, and Su-In Lee. From local explanations to global understanding with explainable AI for trees. *Nature Machine Intelligence*, 2(1):56–67, 2020. doi:10.1038/s42256-019-0138-9.
- [38] Scott M. Lundberg and Su-In Lee. A unified approach to interpreting model predictions. In *Advances in Neural Information Processing Systems*, volume 30, pages 4766–4777, 2017.
- [39] Michael Manville. Unpermitted: The history and economics of land-use regulation. *Journal of the American Planning Association*, 82(1):100–102, 2016. doi:10.1080/01944363.2015.1111666.
- [40] Felicity Maxwell. Executive director, texans for housing. semi-structured interview regarding protest tactics and the legislative impact of hb 24. Personal communication, February 2026.
- [41] Brian J. McCabe. *No Place Like Home: Wealth, Community, and the Politics of Homeownership*. Oxford University Press, 2016. doi:10.1093/acprof:oso/9780190270469.001.0001.
- [42] Justin McCrary. Manipulation of the running variable in the regression discontinuity design: A density test. *Journal of Econometrics*, 142(2):698–714, 2008. doi:10.1016/j.jeconom.2007.05.005.
- [43] Margaret Mitchell, Simone Wu, Andrew Zaldivar, Parker Barnes, Lucy Vasserman, Ben Hutchinson, Elena Spitzer, Inioluwa Deborah Raji, and Timnit Gebru. Model cards for model reporting. In *Proceedings of the Conference on Fairness, Accountability, and Transparency (FAT* ’19)*, pages 220–229. ACM, 2019. doi:10.1145/3287560.3287596.
- [44] Sendhil Mullainathan and Jann Spiess. Machine learning: An applied econometric approach. *Journal of Economic Perspectives*, 31(2):87–106, 2017. doi:10.1257/jep.31.2.87.
- [45] National Institute of Standards and Technology. Artificial intelligence risk management framework (AI RMF 1.0). Technical Report NIST AI 100-1, U.S. Department of Commerce, 2023. URL: <https://www.nist.gov/artificial-intelligence/executive-order-safe-secure-and-trustworthy-artificial-intelligence>, doi:10.6028/NIST.AI.100-1.

- [46] Ziad Obermeyer, Brian Powers, Christine Vogeli, and Sendhil Mullainathan. Dissecting racial bias in an algorithm used to manage the health of populations. *Science*, 366(6464):447–453, 2019. doi:10.1126/science.aax2342.
- [47] Cathy O’Neil. *Weapons of Math Destruction: How Big Data Increases Inequality and Threatens Democracy*. Crown Publishers, 2016.
- [48] Rolf Pendall. Opposition to housing: NIMBY and beyond. *Urban Affairs Review*, 35(1):112–136, 1999. doi:10.1177/107808749903500106.
- [49] John M. Quigley and Larry A. Rosenthal. The effects of land-use regulation on the price of housing: What do we know? what can we learn? *Cityscape*, 8(1):69–137, 2005.
- [50] Inioluwa Deborah Raji, Andrew Smart, Rebecca N. White, Margaret Mitchell, Timnit Gebru, Ben Hutchinson, Jamila Smith-Loud, Daniel Theron, and Parker Barnes. Closing the AI accountability gap: Defining an end-to-end framework for internal algorithmic auditing. In *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency (FAT* ’20)*, pages 33–44. ACM, 2020. doi:10.1145/3351095.3372873.
- [51] Jason Richardson, Bruce Mitchell, and Juan Franco. Dirty deeds: The past and future of redlining. Technical report, National Community Reinvestment Coalition, 2019. URL: <https://ncrc.org/wp-content/uploads/2019/04/NCRC-Dirty-Deeds-The-Past-And-Future-Of-Redlining.pdf>.
- [52] Jonathan Roth, Pedro H.C. Sant’Anna, Alyssa Bilinski, and John Poe. What’s trending in difference-in-differences? a synthesis of the recent econometrics literature. *Journal of Econometrics*, 235(2):2218–2244, 2023. doi:10.1016/j.jeconom.2023.03.008.
- [53] Carissa Schively. Understanding the NIMBY and LULU phenomena: Reassessing our knowledge base and informing future research. *Journal of Planning Literature*, 21(3):255–266, 2007. doi:10.1177/0885412206295845.
- [54] Andrew D. Selbst, Danah Boyd, Sorelle A. Friedler, Suresh Venkatasubramanian, and Janet Vertesi. Fairness and abstraction in sociotechnical systems. In *Proceedings of the Conference on Fairness, Accountability, and Transparency (FAT* ’19)*, pages 59–68. ACM, 2019. doi:10.1145/3287560.3287598.
- [55] Galit Shmueli. To explain or to predict? *Statistical Science*, 25(3):289–310, 2010. doi:10.1214/10-STS330.
- [56] Gaurav Sood and Suriyan Laohaprapanon. *ethnicolr*: Predict race and ethnicity from name. *CRAN/PyPI*, 2018. URL: <https://github.com/appeler/ethnicolr>.
- [57] Texas Legislature. House bill 24, 89th legislature, regular session, 2025. Effective September 1, 2025. Codified as Texas Local Government Code §211.0061. URL: <https://capitol.texas.gov/BillLookup/History.aspx?LegSess=89R&Bill=HB24>.
- [58] Texas State Legislature. Texas local government code § 211.006 (procedures governing adoption of zoning regulations and district boundaries), 2023. URL: <https://statutes.capitol.texas.gov/Docs/LG/htm/LG.211.htm>.
- [59] Texas State Legislature. Texas local government code § 211.0061 (protests and supermajority voting), 2025. URL: <https://statutes.capitol.texas.gov/Docs/LG/htm/LG.211.htm>.

- [60] Jonathan Tomko. Principal planner, city of austin. semi-structured interview regarding predictive tools and administrative adoption. Personal communication, February 2026.
- [61] Travis Central Appraisal District. Electronic appraisal roll submission (EARS), 2018–2025. Texas Comptroller of Public Accounts Property Tax Data Portal, 2025. URL: <https://comptroller.texas.gov/taxes/property-tax/>.
- [62] Travis County Clerk. November 2022 election results. Technical report, Travis County, TX, 2022. URL: <https://www.traviscountyclerk.org/eclerk/content/election-results>.
- [63] U.S. Census Bureau. American community survey 5-year estimates (2009–2023). Technical report, U.S. Department of Commerce, 2023. URL: <https://data.census.gov>.
- [64] Heather K. Way. *The Racist Roots of Austin’s Segregation*. University of Texas Press, 2018.
- [65] Andrew H. Whittemore. Racial and class bias in exclusionary zoning: The case of Maywood, Illinois. *Journal of Planning History*, 13(3):250–272, 2014. doi:10.1177/1538513213511195.

Appendix Guide. The appendix is organized into three blocks. *Qualitative Documentation* (Appendix A) contains recruitment, IRB, interview, petition, and panel artifacts mapping to Chapter 3. *Quantitative Methods and Robustness* (Appendix B) contains detailed modeling configurations, variable mapping, calibration transparency, and exploratory prediction diagnostics mapping to Chapters 4 and 5. *Policy Event Studies* (Appendix C) contains independent tracking of the HOME Initiative shock. Detailed navigation is provided by the table of contents.

A Qualitative Documentation

A.1 Interview Recruitment Overview

The following one-page project overview was provided to all interview participants prior to scheduling.

Research Project Overview

Predicting NIMBYism in Austin's Housing Reform Era

Project Abstract

This research forecasts opposition to development using AI and interprets the mechanics of that opposition in Austin, Texas. By analyzing petition protest letters, public hearing transcripts, and zoning application data, the study evaluates the democratic implications of using predictive analytics to anticipate political behavior. The project employs predictive modeling, natural language processing, and qualitative analysis to predict opposition behavior, identify recurring patterns in opposition rhetoric, and examine their association with planning outcomes.

Study Objectives

Identify Patterns: Characterize the opposition rhetoric and arguments used in protest petitions to oppose rezoning and identify demographic predictors of this opposition.

Evaluate Predictive Tools: Assess whether administrative data and predictive models can accurately forecast opposition to specific development projects.

Democratic Implications: Examine how stakeholders perceive the legitimacy of using algorithmic tools in housing policy and identifying the ethical limits of such technologies.

Participant Role

We are seeking perspectives from urban planners, developers, neighborhood association leaders, and community organizers.

Activity: A 45-60 minute one-on-one interview conducted virtually or in person.

Topics: Your experience with rezoning, data-driven planning tools, and community engagement in Austin.

Data Privacy and Ethical Standards

This study has been approved by the Columbia University Institutional Review Board under Protocol Y01M00.

Participation is entirely voluntary.

You may decline to answer any specific question or discontinue the interview at any time.

Identities will be kept strictly confidential in resulting publications unless you grant explicit permission to be named.

Researcher Contact

Daniel Hardesty Lewis
Graduate Researcher
Email: dl3645@columbia.edu

IRB Contact

Columbia University Morningside IRB
Phone: (212) 851-7040
Email: AskIRB@columbia.edu

A.2 IRB Protocol Summary

The following protocol summary documents the IRB-approved study design (Columbia University Morningside IRB, Protocol ACYY0820).

COLUMBIA UNIVERSITY
Graduate School of Architecture, Planning and Preservation

Protocol Summary
ACYY0820

Study Protocol Summary

Predicting NIMBYism in Austin's Housing Reform Era

Key Personnel

Principal Investigator (PI): Hiba Bou Akar, Ph.D.

Associate Professor, Urban Planning

Columbia University GSAPP

Email: hb2541@columbia.edu

Student Researcher: Daniel Lewis

M.S. Urban Planning Student, Columbia University GSAPP

Email: dl3645@columbia.edu

IRB Contact

If you have questions about your rights or welfare as a research participant, you may contact:

Columbia University Morningside Institutional Review Board (IRB)

Phone: (212) 851-7040

Email: AskIRB@columbia.edu

Study Purpose

Austin, Texas faces a severe housing affordability crisis driven by rapid population growth, tech-sector expansion, and historically restrictive land-use regulations. Recent housing reforms have coincided with organized neighborhood opposition to development using tools such as protest petitions, public testimony, and legal challenges. This study aims to:

1. Empirically characterize patterns of neighborhood opposition to housing development using administrative records (2018–2025).
2. Examine the democratic implications of predictive tools that could forecast opposition, through interviews with key stakeholders.

Research Questions

Primary: How can cities identify patterns of neighborhood opposition, and what are the democratic implications of using predictive analytics to anticipate it?

Sub-questions:

- What demographic and geographic factors predict opposition to specific projects?
- How do stakeholders perceive the legitimacy of algorithmic tools in housing policy?
- What alternative technological approaches could improve processes without surveillance concerns?

Study Design

This mixed-methods, minimal-risk study combines:

- **Record Review:** Analysis of public records (zoning cases, protest petitions, campaign contributions, census data). De-identified datasets will use coded IDs.
- **Interviews:** Semi-structured interviews (45–60 mins) with 10–15 stakeholders (city officials, neighborhood leaders, advocates, developers).
- **Observation:** Review of public meetings (Council, commissions).

Participation Details

Interviews: Conducted in-person or via secure videoconference. With permission, they are audio-recorded for transcription. Participation is voluntary; you may decline to answer or stop at any time.

Risks: Minimal risk. Potential breach of confidentiality is mitigated by secure storage, encryption, and reporting only aggregate/de-identified results.

Benefits: No direct benefit. Indirect benefits include informing more equitable planning practices.

A.3 Semi-Structured Interview Guide

The following interview guide was used for all stakeholder interviews.

Predicting NIMBYism in Austin's Housing Reform Era

Purpose

To explore how key stakeholders understand neighborhood opposition to housing development in Austin and how they view the potential use of predictive or algorithmic tools in planning and zoning.

Opening

1. Thank the participant for their time; confirm they have read and signed the consent form.
2. Confirm whether they consent to audio recording.
3. Remind them they may skip any question or stop at any time.

Part 1: Background and Role

1. Can you briefly describe your current role and how it relates to housing, planning, or land use in Austin?
2. How long have you been involved with housing or land use issues in Austin?
3. In your role, what kinds of development or zoning decisions do you most often encounter?

Part 2: Experiences with Neighborhood Opposition

1. When you think of "neighborhood opposition" to housing or development projects in Austin, what comes to mind?
2. Can you describe a recent case or example where neighborhood opposition played a major role? What happened?
3. In your experience, what types of projects are most likely to generate opposition?
4. Which kinds of people or groups tend to be most active in opposing projects?

Part 3: Perceptions of Fairness and Participation

1. Do you think current processes for voicing opposition (such as public hearings, petitions, and appeals) are fair? Why or why not?
2. Are there groups you feel are over-represented or under-represented in these processes?
3. How do you think opposition affects whose interests are ultimately reflected in housing and zoning decisions?

Part 4: Views on Predictive and Algorithmic Tools

1. Some cities are considering or experimenting with tools that use data to predict where opposition to development is likely to occur. Have you encountered ideas like this before?
2. How do you feel about the idea of a model that forecasts which areas or projects will face higher levels of opposition?
3. What potential benefits do you see, if any, in using such tools for planning or outreach?
4. What potential risks or harms concern you most? For example, concerns about profiling, surveillance, or chilling participation.
5. Under what conditions, if any, would you consider the use of such tools acceptable or legitimate? (e.g., transparency, public input, limits on use).

B Quantitative Methods and Robustness

B.1 Expanded Methods Appendix: Technical Details

B.1.1 Hyperparameter Search Grids

The unified architectural benchmark utilized the following randomized search bounds in 5-fold temporal cross-validation:

- **Tree Ensembles (CatBoost):** Learning Rate $\sim \log \text{Uniform}[0.01, 0.3]$, Max Depth $\in \{4, 6, 8, 10\}$.
- **Deep Architectures (MLP/TabNet):** Layers $\in \{[32], [64, 32], [128, 64, 32]\}$, Weight Decay $\sim \log \text{Uniform}[1e-5, 1e-2]$.
- **Regularized Linear (Logistic L2):** C (Inverse Regularization) $\sim \log \text{Uniform}[0.01, 10]$.
- **Distributionally Robust (Anchor Regression):** Alpha $\sim \text{Uniform}[1, 100]$.

All algorithms were tuned exclusively to maximize Out-Of-Fold Precision-Recall AUC (PR-AUC) under strict temporal separation.




Figure 14: **Out-of-distribution optimization surfaces for benchmark architectures.** Empirically computed grid-search PR-AUC heatmaps across the four formalized model families (Tree, Deep, Linear, and Causal). Architectures like Deep Tabular overfit catastrophically under high capacity without massive L2 constraint, while Tree Ensembles plateau gracefully.

B.2 Temporal Drift Threshold Sensitivity

This appendix replicates the primary temporal drift and longitudinal stability exhibits across varying petition percentage thresholds ($> 0\%$, $> 5\%$, $> 10\%$, $> 15\%$, $> 20\%$, and $> 25\%$) to validate whether architectural dependence patterns isolate to the statutory cutoff or represent generalized predictive dynamics at all boundaries.

Algorithmic Predictability and Statutory Mobilization Constraints. Figure 15 visualizes the performance trajectory of the dominant architecture (Max-of-Family PR-AUC) across these six thresholds. The algorithmic sensitivity indicates a highly volatile social mechanism governing neighbor mobilization. Predictive signal does not stabilize or peak consistently around the statutory cutoff; instead, performance fluctuates erratically across anchors and thresholds. This volatility suggests that formal petition generation is not a deterministic institutional barricade, but rather a contingent social phenomenon that administrative records cannot perfectly delineate.




Figure 15: Algorithmic predictability (out-of-distribution PR-AUC) of the dominant architectural family across progressive petition area thresholds for all rolling origins. Performance fluctuates significantly across years and thresholds, indicating that the 'organized' capacity of a neighborhood is a contingent social process rather than a static response to the legal threshold.

B.2.1 Threshold: $>0\%$

Table 11: Temporal predictive drift: PR-AUC decay by algorithm.

Model	Anchor Training	2018	2019	2020	2021	2022	2023	2024
CatBoost	Pre-2018	0.947	0.705	0.646	0.522	0.246	0.945	0.659
CatBoost	Pre-2019	—	0.645	0.697	0.630	0.296	0.863	0.605
CatBoost	Pre-2020	—	—	0.640	0.656	0.219	0.823	0.593
CatBoost	Pre-2021	—	—	—	0.755	0.220	1.000	0.761
CatBoost	Pre-2022	—	—	—	—	0.245	0.895	0.609
CatBoost	Pre-2023	—	—	—	—	—	0.804	0.444
CatBoost	Pre-2024	—	—	—	—	—	—	0.520
XGBoost	Pre-2018	0.912	0.849	0.617	0.537	0.277	0.948	0.634
XGBoost	Pre-2019	—	0.831	0.607	0.588	0.249	0.904	0.639
XGBoost	Pre-2020	—	—	0.637	0.640	0.304	0.941	0.559
XGBoost	Pre-2021	—	—	—	0.601	0.222	0.956	0.664
XGBoost	Pre-2022	—	—	—	—	0.330	0.942	0.525
XGBoost	Pre-2023	—	—	—	—	—	0.785	0.458
XGBoost	Pre-2024	—	—	—	—	—	—	0.420
Logistic (L2)	Pre-2018	0.947	0.950	0.552	0.521	0.260	0.741	0.544
Logistic (L2)	Pre-2019	—	0.950	0.567	0.526	0.246	0.764	0.533
Logistic (L2)	Pre-2020	—	—	0.549	0.519	0.251	0.749	0.544
Logistic (L2)	Pre-2021	—	—	—	0.521	0.251	0.714	0.533
Logistic (L2)	Pre-2022	—	—	—	—	0.277	0.728	0.500
Logistic (L2)	Pre-2023	—	—	—	—	—	0.730	0.340
Logistic (L2)	Pre-2024	—	—	—	—	—	—	0.362
Spatial-FE Logistic	Pre-2018	0.942	0.839	0.651	0.452	0.235	0.595	0.730
Spatial-FE Logistic	Pre-2019	—	0.849	0.709	0.488	0.224	0.641	0.511
Spatial-FE Logistic	Pre-2020	—	—	0.672	0.552	0.220	0.720	0.564
Spatial-FE Logistic	Pre-2021	—	—	—	0.532	0.206	0.728	0.559
Spatial-FE Logistic	Pre-2022	—	—	—	—	0.305	0.724	0.511
Spatial-FE Logistic	Pre-2023	—	—	—	—	—	0.698	0.339
Spatial-FE Logistic	Pre-2024	—	—	—	—	—	—	0.325
TabNet	Pre-2018	0.898	0.934	0.743	0.404	0.387	0.744	0.452
TabNet	Pre-2019	—	0.714	0.528	0.375	0.147	0.184	0.128
TabNet	Pre-2020	—	—	0.638	0.419	0.194	0.410	0.214
TabNet	Pre-2021	—	—	—	0.615	0.196	0.372	0.239
TabNet	Pre-2022	—	—	—	—	0.210	0.469	0.489
TabNet	Pre-2023	—	—	—	—	—	0.641	0.322
TabNet	Pre-2024	—	—	—	—	—	—	0.441
Anchor Regression (Causal)	Pre-2018	0.825	0.788	0.671	0.519	0.251	0.682	0.510
Anchor Regression (Causal)	Pre-2019	—	0.318	0.520	0.432	0.200	0.676	0.625
Anchor Regression (Causal)	Pre-2020	—	—	0.753	0.385	0.216	0.513	0.534
Anchor Regression (Causal)	Pre-2021	—	—	—	0.643	0.251	0.785	0.473
Anchor Regression (Causal)	Pre-2022	—	—	—	—	0.219	0.650	0.492
Anchor Regression (Causal)	Pre-2023	—	—	—	—	—	0.302	0.331
Anchor Regression (Causal)	Pre-2024	—	—	—	—	—	—	0.462

Table 12: Temporal predictive drift: PR-AUC lift by algorithm.

Model	Anchor Training	2018	2019	2020	2021	2022	2023	2024
CatBoost	Pre-2018	6.218	5.919	8.167	7.977	10.611	18.998	28.349
CatBoost	Pre-2019	—	5.417	8.814	9.630	12.771	17.346	26.032
CatBoost	Pre-2020	—	—	8.097	10.020	9.448	16.538	25.483
CatBoost	Pre-2021	—	—	—	11.545	9.510	20.100	32.728
CatBoost	Pre-2022	—	—	—	—	10.591	17.998	26.175
CatBoost	Pre-2023	—	—	—	—	—	16.157	19.104
CatBoost	Pre-2024	—	—	—	—	—	—	22.360
XGBoost	Pre-2018	5.986	7.135	7.795	8.206	11.950	19.065	27.274
XGBoost	Pre-2019	—	6.976	7.675	8.985	10.770	18.164	27.489
XGBoost	Pre-2020	—	—	8.051	9.790	13.147	18.907	24.049
XGBoost	Pre-2021	—	—	—	9.186	9.604	19.212	28.571
XGBoost	Pre-2022	—	—	—	—	14.254	18.937	22.592
XGBoost	Pre-2023	—	—	—	—	—	15.774	19.691
XGBoost	Pre-2024	—	—	—	—	—	—	18.050
Logistic (L2)	Pre-2018	6.220	7.976	6.975	7.962	11.220	14.896	23.402
Logistic (L2)	Pre-2019	—	7.981	7.165	8.035	10.612	15.366	22.933
Logistic (L2)	Pre-2020	—	—	6.940	7.927	10.859	15.064	23.411
Logistic (L2)	Pre-2021	—	—	—	7.957	10.850	14.361	22.933
Logistic (L2)	Pre-2022	—	—	—	—	11.969	14.629	21.500
Logistic (L2)	Pre-2023	—	—	—	—	—	14.670	14.630
Logistic (L2)	Pre-2024	—	—	—	—	—	—	15.582
Spatial-FE Logistic	Pre-2018	6.187	7.049	8.227	6.903	10.132	11.950	31.370
Spatial-FE Logistic	Pre-2019	—	7.134	8.968	7.463	9.672	12.883	21.978
Spatial-FE Logistic	Pre-2020	—	—	8.495	8.438	9.519	14.475	24.271
Spatial-FE Logistic	Pre-2021	—	—	—	8.131	8.884	14.636	24.025
Spatial-FE Logistic	Pre-2022	—	—	—	—	13.182	14.554	21.978
Spatial-FE Logistic	Pre-2023	—	—	—	—	—	14.039	14.593
Spatial-FE Logistic	Pre-2024	—	—	—	—	—	—	13.985
TabNet	Pre-2018	5.895	7.842	9.389	6.182	16.708	14.954	19.448
TabNet	Pre-2019	—	5.995	6.672	5.735	6.361	3.694	5.522
TabNet	Pre-2020	—	—	8.067	6.402	8.387	8.244	9.222
TabNet	Pre-2021	—	—	—	9.401	8.451	7.471	10.257
TabNet	Pre-2022	—	—	—	—	9.059	9.427	21.022
TabNet	Pre-2023	—	—	—	—	—	12.890	13.849
TabNet	Pre-2024	—	—	—	—	—	—	18.968
Anchor Regression (Causal)	Pre-2018	5.418	6.621	8.489	7.926	10.834	13.705	21.920
Anchor Regression (Causal)	Pre-2019	—	2.670	6.571	6.603	8.650	13.582	26.875
Anchor Regression (Causal)	Pre-2020	—	—	9.520	5.880	9.334	10.304	22.975
Anchor Regression (Causal)	Pre-2021	—	—	—	9.833	10.846	15.784	20.344
Anchor Regression (Causal)	Pre-2022	—	—	—	—	9.457	13.065	21.159
Anchor Regression (Causal)	Pre-2023	—	—	—	—	—	6.072	14.229
Anchor Regression (Causal)	Pre-2024	—	—	—	—	—	—	19.873

Table 13: **Max-of-Family dominance.**

Anchor Training	2018	2019	2020	2021	2022	2023	2024
Panel A: Maximum Absolute PR-AUC							
Pre-2018	0.947 (Linear)	0.950 (Linear)	0.743 (Deep)	0.537 (Trees)	0.387 (Deep)	0.948 (Trees)	0.730 (Linear)
Pre-2019	—	0.950 (Linear)	0.709 (Linear)	0.630 (Trees)	0.296 (Trees)	0.904 (Trees)	0.639 (Trees)
Pre-2020	—	—	0.753 (Causal)	0.656 (Trees)	0.304 (Trees)	0.941 (Trees)	0.593 (Trees)
Pre-2021	—	—	—	0.755 (Trees)	0.251 (Linear)	1.000 (Trees)	0.761 (Trees)
Pre-2022	—	—	—	—	0.330 (Trees)	0.942 (Trees)	0.609 (Trees)
Pre-2023	—	—	—	—	—	0.804 (Trees)	0.458 (Trees)
Pre-2024	—	—	—	—	—	—	0.520 (Trees)
Panel B: Maximum Relative PR-AUC Lift							
Pre-2018	6.22 (Linear)	7.98 (Linear)	9.39 (Deep)	8.21 (Trees)	16.71 (Deep)	19.06 (Trees)	31.37 (Linear)
Pre-2019	—	7.98 (Linear)	8.97 (Linear)	9.63 (Trees)	12.77 (Trees)	18.16 (Trees)	27.49 (Trees)
Pre-2020	—	—	9.52 (Causal)	10.02 (Trees)	13.15 (Trees)	18.91 (Trees)	25.48 (Trees)
Pre-2021	—	—	—	11.55 (Trees)	10.85 (Linear)	20.10 (Trees)	32.73 (Trees)
Pre-2022	—	—	—	—	14.25 (Trees)	18.94 (Trees)	26.18 (Trees)
Pre-2023	—	—	—	—	—	16.16 (Trees)	19.69 (Trees)
Pre-2024	—	—	—	—	—	—	22.36 (Trees)

B.2.2 Threshold: $>5\%$

Table 14: Temporal predictive drift: PR-AUC decay by algorithm.

Model	Anchor Training	2018	2019	2020	2021	2022	2023	2024
CatBoost	Pre-2018	0.956	0.912	0.568	0.564	0.313	0.832	0.875
CatBoost	Pre-2019	—	0.762	0.622	0.666	0.244	0.770	0.854
CatBoost	Pre-2020	—	—	0.507	0.624	0.305	0.831	0.732
CatBoost	Pre-2021	—	—	—	0.716	0.355	0.983	0.568
CatBoost	Pre-2022	—	—	—	—	0.684	0.950	0.385
CatBoost	Pre-2023	—	—	—	—	—	0.970	0.352
CatBoost	Pre-2024	—	—	—	—	—	—	0.568
XGBoost	Pre-2018	0.923	0.987	0.601	0.511	0.249	0.917	0.625
XGBoost	Pre-2019	—	0.915	0.583	0.520	0.244	0.895	0.668
XGBoost	Pre-2020	—	—	0.571	0.519	0.237	0.970	0.649
XGBoost	Pre-2021	—	—	—	0.635	0.242	0.912	0.650
XGBoost	Pre-2022	—	—	—	—	0.301	0.970	0.568
XGBoost	Pre-2023	—	—	—	—	—	0.909	0.390
XGBoost	Pre-2024	—	—	—	—	—	—	0.575
Logistic (L2)	Pre-2018	0.980	0.987	0.453	0.368	0.264	0.587	0.392
Logistic (L2)	Pre-2019	—	0.987	0.459	0.385	0.263	0.599	0.392
Logistic (L2)	Pre-2020	—	—	0.447	0.382	0.288	0.599	0.415
Logistic (L2)	Pre-2021	—	—	—	0.376	0.308	0.665	0.392
Logistic (L2)	Pre-2022	—	—	—	—	0.317	0.585	0.413
Logistic (L2)	Pre-2023	—	—	—	—	—	0.594	0.390
Logistic (L2)	Pre-2024	—	—	—	—	—	—	0.415
Spatial-FE Logistic	Pre-2018	0.998	0.959	0.629	0.439	0.268	0.555	0.548
Spatial-FE Logistic	Pre-2019	—	0.962	0.631	0.433	0.247	0.557	0.531
Spatial-FE Logistic	Pre-2020	—	—	0.568	0.394	0.258	0.562	0.540
Spatial-FE Logistic	Pre-2021	—	—	—	0.316	0.257	0.608	0.540
Spatial-FE Logistic	Pre-2022	—	—	—	—	0.339	0.613	0.540
Spatial-FE Logistic	Pre-2023	—	—	—	—	—	0.602	0.526
Spatial-FE Logistic	Pre-2024	—	—	—	—	—	—	0.552
TabNet	Pre-2018	0.939	0.613	0.426	0.364	0.117	0.154	0.090
TabNet	Pre-2019	—	0.754	0.790	0.371	0.218	0.534	0.282
TabNet	Pre-2020	—	—	0.595	0.424	0.443	0.708	0.582
TabNet	Pre-2021	—	—	—	0.594	0.378	0.722	0.338
TabNet	Pre-2022	—	—	—	—	0.231	0.746	0.608
TabNet	Pre-2023	—	—	—	—	—	0.549	0.565
TabNet	Pre-2024	—	—	—	—	—	—	0.510
Anchor Regression (Causal)	Pre-2018	0.820	0.832	0.421	0.328	0.281	0.636	0.531
Anchor Regression (Causal)	Pre-2019	—	0.582	0.542	0.339	0.294	0.604	0.436
Anchor Regression (Causal)	Pre-2020	—	—	0.525	0.380	0.265	0.662	0.392
Anchor Regression (Causal)	Pre-2021	—	—	—	0.409	0.417	0.605	0.514
Anchor Regression (Causal)	Pre-2022	—	—	—	—	0.350	0.506	0.261
Anchor Regression (Causal)	Pre-2023	—	—	—	—	—	0.696	0.500
Anchor Regression (Causal)	Pre-2024	—	—	—	—	—	—	0.472

Table 15: **Temporal predictive drift: PR-AUC lift by algorithm.**

Model	Anchor Training	2018	2019	2020	2021	2022	2023	2024
CatBoost	Pre-2018	6.278	7.660	8.382	10.058	16.907	16.727	47.031
CatBoost	Pre-2019	—	6.397	9.175	11.871	13.164	15.484	45.911
CatBoost	Pre-2020	—	—	7.484	11.124	16.459	16.696	39.353
CatBoost	Pre-2021	—	—	—	12.765	19.157	19.765	30.544
CatBoost	Pre-2022	—	—	—	—	36.940	19.089	20.689
CatBoost	Pre-2023	—	—	—	—	—	19.493	18.898
CatBoost	Pre-2024	—	—	—	—	—	—	30.544
XGBoost	Pre-2018	6.061	8.292	8.872	9.117	13.432	18.422	33.594
XGBoost	Pre-2019	—	7.687	8.593	9.276	13.154	17.998	35.897
XGBoost	Pre-2020	—	—	8.429	9.255	12.779	19.493	34.874
XGBoost	Pre-2021	—	—	—	11.325	13.048	18.334	34.938
XGBoost	Pre-2022	—	—	—	—	16.232	19.493	30.544
XGBoost	Pre-2023	—	—	—	—	—	18.261	20.988
XGBoost	Pre-2024	—	—	—	—	—	—	30.906
Logistic (L2)	Pre-2018	6.436	8.292	6.689	6.556	14.250	11.798	21.052
Logistic (L2)	Pre-2019	—	8.292	6.764	6.862	14.175	12.036	21.052
Logistic (L2)	Pre-2020	—	—	6.595	6.812	15.557	12.036	22.332
Logistic (L2)	Pre-2021	—	—	—	6.713	16.609	13.376	21.052
Logistic (L2)	Pre-2022	—	—	—	—	17.132	11.752	22.209
Logistic (L2)	Pre-2023	—	—	—	—	—	11.940	20.972
Logistic (L2)	Pre-2024	—	—	—	—	—	—	22.332
Spatial-FE Logistic	Pre-2018	6.553	8.058	9.282	7.825	14.483	11.147	29.458
Spatial-FE Logistic	Pre-2019	—	8.079	9.309	7.725	13.358	11.205	28.562
Spatial-FE Logistic	Pre-2020	—	—	8.378	7.026	13.911	11.305	29.051
Spatial-FE Logistic	Pre-2021	—	—	—	5.627	13.904	12.228	29.051
Spatial-FE Logistic	Pre-2022	—	—	—	—	18.300	12.311	29.051
Spatial-FE Logistic	Pre-2023	—	—	—	—	—	12.091	28.288
Spatial-FE Logistic	Pre-2024	—	—	—	—	—	—	29.648
TabNet	Pre-2018	6.164	5.147	6.288	6.493	6.335	3.099	4.826
TabNet	Pre-2019	—	6.334	11.652	6.621	11.794	10.729	15.143
TabNet	Pre-2020	—	—	8.778	7.563	23.936	14.222	31.290
TabNet	Pre-2021	—	—	—	10.585	20.411	14.522	18.173
TabNet	Pre-2022	—	—	—	—	12.482	14.993	32.698
TabNet	Pre-2023	—	—	—	—	—	11.029	30.394
TabNet	Pre-2024	—	—	—	—	—	—	27.408
Anchor Regression (Causal)	Pre-2018	5.388	6.991	6.215	5.853	15.199	12.783	28.539
Anchor Regression (Causal)	Pre-2019	—	4.893	7.994	6.038	15.854	12.147	23.441
Anchor Regression (Causal)	Pre-2020	—	—	7.748	6.780	14.308	13.316	21.089
Anchor Regression (Causal)	Pre-2021	—	—	—	7.286	22.500	12.162	27.622
Anchor Regression (Causal)	Pre-2022	—	—	—	—	18.900	10.174	14.023
Anchor Regression (Causal)	Pre-2023	—	—	—	—	—	13.991	26.875
Anchor Regression (Causal)	Pre-2024	—	—	—	—	—	—	25.382

Table 16: **Max-of-Family dominance.**

Anchor Training	2018	2019	2020	2021	2022	2023	2024
Panel A: Maximum Absolute PR-AUC							
Pre-2018	0.998 (Linear)	0.987 (Trees)	0.629 (Linear)	0.564 (Trees)	0.313 (Trees)	0.917 (Trees)	0.875 (Trees)
Pre-2019	—	0.987 (Linear)	0.790 (Deep)	0.666 (Trees)	0.294 (Causal)	0.895 (Trees)	0.854 (Trees)
Pre-2020	—	—	0.595 (Deep)	0.624 (Trees)	0.443 (Deep)	0.970 (Trees)	0.732 (Trees)
Pre-2021	—	—	—	0.716 (Trees)	0.417 (Causal)	0.983 (Trees)	0.650 (Trees)
Pre-2022	—	—	—	—	0.684 (Trees)	0.970 (Trees)	0.608 (Deep)
Pre-2023	—	—	—	—	—	0.970 (Trees)	0.565 (Deep)
Pre-2024	—	—	—	—	—	—	0.575 (Trees)
Panel B: Maximum Relative PR-AUC Lift							
Pre-2018	6.55 (Linear)	8.29 (Linear)	9.28 (Linear)	10.06 (Trees)	16.91 (Trees)	18.42 (Trees)	47.03 (Trees)
Pre-2019	—	8.29 (Linear)	11.65 (Deep)	11.87 (Trees)	15.85 (Causal)	18.00 (Trees)	45.91 (Trees)
Pre-2020	—	—	8.78 (Deep)	11.12 (Trees)	23.94 (Deep)	19.49 (Trees)	39.35 (Trees)
Pre-2021	—	—	—	12.76 (Trees)	22.50 (Causal)	19.76 (Trees)	34.94 (Trees)
Pre-2022	—	—	—	—	36.94 (Trees)	19.49 (Trees)	32.70 (Deep)
Pre-2023	—	—	—	—	—	19.49 (Trees)	30.39 (Deep)
Pre-2024	—	—	—	—	—	—	30.91 (Trees)

B.2.3 Threshold: $>10\%$

Table 17: Temporal predictive drift: PR-AUC decay by algorithm.

Model	Anchor Training	2018	2019	2020	2021	2022	2023	2024
CatBoost	Pre-2018	0.929	0.524	0.560	0.416	0.317	0.594	0.622
CatBoost	Pre-2019	—	0.494	0.605	0.440	0.330	0.728	0.580
CatBoost	Pre-2020	—	—	0.510	0.385	0.304	0.524	0.367
CatBoost	Pre-2021	—	—	—	0.682	0.500	0.676	0.339
CatBoost	Pre-2022	—	—	—	—	0.607	0.689	0.339
CatBoost	Pre-2023	—	—	—	—	—	0.709	0.360
CatBoost	Pre-2024	—	—	—	—	—	—	0.360
XGBoost	Pre-2018	0.909	0.609	0.499	0.419	0.582	0.597	0.565
XGBoost	Pre-2019	—	0.461	0.560	0.464	0.317	0.583	0.649
XGBoost	Pre-2020	—	—	0.642	0.436	0.424	0.479	0.338
XGBoost	Pre-2021	—	—	—	0.432	0.320	0.518	0.649
XGBoost	Pre-2022	—	—	—	—	0.424	0.628	0.371
XGBoost	Pre-2023	—	—	—	—	—	0.609	0.419
XGBoost	Pre-2024	—	—	—	—	—	—	0.421
Logistic (L2)	Pre-2018	0.952	0.687	0.453	0.374	0.330	0.453	0.399
Logistic (L2)	Pre-2019	—	0.654	0.466	0.402	0.330	0.519	0.365
Logistic (L2)	Pre-2020	—	—	0.491	0.327	0.280	0.463	0.395
Logistic (L2)	Pre-2021	—	—	—	0.382	0.306	0.484	0.364
Logistic (L2)	Pre-2022	—	—	—	—	0.412	0.472	0.365
Logistic (L2)	Pre-2023	—	—	—	—	—	0.470	0.360
Logistic (L2)	Pre-2024	—	—	—	—	—	—	0.340
Spatial-FE Logistic	Pre-2018	0.802	0.734	0.510	0.330	0.244	0.452	0.667
Spatial-FE Logistic	Pre-2019	—	0.704	0.578	0.319	0.213	0.498	0.333
Spatial-FE Logistic	Pre-2020	—	—	0.565	0.290	0.221	0.530	0.399
Spatial-FE Logistic	Pre-2021	—	—	—	0.311	0.223	0.558	0.374
Spatial-FE Logistic	Pre-2022	—	—	—	—	0.225	0.462	0.349
Spatial-FE Logistic	Pre-2023	—	—	—	—	—	0.443	0.380
Spatial-FE Logistic	Pre-2024	—	—	—	—	—	—	0.360
TabNet	Pre-2018	0.943	0.558	0.604	0.446	0.217	0.452	0.317
TabNet	Pre-2019	—	0.874	0.536	0.353	0.250	0.732	0.380
TabNet	Pre-2020	—	—	0.618	0.519	0.127	0.655	0.155
TabNet	Pre-2021	—	—	—	0.365	0.264	0.501	0.232
TabNet	Pre-2022	—	—	—	—	0.066	0.070	0.041
TabNet	Pre-2023	—	—	—	—	—	0.413	0.364
TabNet	Pre-2024	—	—	—	—	—	—	0.378
Anchor Regression (Causal)	Pre-2018	0.986	0.398	0.191	0.473	0.152	0.177	0.063
Anchor Regression (Causal)	Pre-2019	—	0.586	0.675	0.362	0.392	0.394	0.296
Anchor Regression (Causal)	Pre-2020	—	—	0.472	0.322	0.329	0.470	0.368
Anchor Regression (Causal)	Pre-2021	—	—	—	0.451	0.288	0.346	0.554
Anchor Regression (Causal)	Pre-2022	—	—	—	—	0.281	0.243	0.572
Anchor Regression (Causal)	Pre-2023	—	—	—	—	—	0.446	0.455
Anchor Regression (Causal)	Pre-2024	—	—	—	—	—	—	0.400

Table 18: **Temporal predictive drift: PR-AUC lift by algorithm.**

Model	Anchor Training	2018	2019	2020	2021	2022	2023	2024
CatBoost	Pre-2018	6.103	6.468	8.263	7.415	17.140	17.063	33.407
CatBoost	Pre-2019	—	6.103	8.925	7.849	17.807	20.896	31.168
CatBoost	Pre-2020	—	—	7.527	6.866	16.416	15.039	19.746
CatBoost	Pre-2021	—	—	—	12.165	27.000	19.416	18.210
CatBoost	Pre-2022	—	—	—	—	32.786	19.794	18.210
CatBoost	Pre-2023	—	—	—	—	—	20.350	19.330
CatBoost	Pre-2024	—	—	—	—	—	—	19.330
XGBoost	Pre-2018	5.971	7.525	7.364	7.471	31.436	17.136	30.394
XGBoost	Pre-2019	—	5.695	8.262	8.280	17.091	16.745	34.874
XGBoost	Pre-2020	—	—	9.470	7.782	22.909	13.766	18.178
XGBoost	Pre-2021	—	—	—	7.712	17.284	14.880	34.874
XGBoost	Pre-2022	—	—	—	—	22.909	18.039	19.924
XGBoost	Pre-2023	—	—	—	—	—	17.488	22.545
XGBoost	Pre-2024	—	—	—	—	—	—	22.652
Logistic (L2)	Pre-2018	6.252	8.489	6.686	6.667	17.847	13.014	21.436
Logistic (L2)	Pre-2019	—	8.079	6.874	7.164	17.815	14.910	19.644
Logistic (L2)	Pre-2020	—	—	7.235	5.829	15.107	13.307	21.212
Logistic (L2)	Pre-2021	—	—	—	6.805	16.529	13.885	19.570
Logistic (L2)	Pre-2022	—	—	—	—	22.275	13.559	19.644
Logistic (L2)	Pre-2023	—	—	—	—	—	13.493	19.372
Logistic (L2)	Pre-2024	—	—	—	—	—	—	18.301
Spatial-FE Logistic	Pre-2018	5.265	9.072	7.529	5.878	13.154	12.983	35.833
Spatial-FE Logistic	Pre-2019	—	8.697	8.532	5.696	11.483	14.309	17.917
Spatial-FE Logistic	Pre-2020	—	—	8.340	5.178	11.916	15.205	21.436
Spatial-FE Logistic	Pre-2021	—	—	—	5.540	12.036	16.012	20.076
Spatial-FE Logistic	Pre-2022	—	—	—	—	12.157	13.277	18.733
Spatial-FE Logistic	Pre-2023	—	—	—	—	—	12.720	20.412
Spatial-FE Logistic	Pre-2024	—	—	—	—	—	—	19.330
TabNet	Pre-2018	6.191	6.895	8.910	7.958	11.700	12.989	17.053
TabNet	Pre-2019	—	10.791	7.910	6.299	13.479	21.025	20.412
TabNet	Pre-2020	—	—	9.120	9.250	6.843	18.807	8.339
TabNet	Pre-2021	—	—	—	6.501	14.279	14.400	12.456
TabNet	Pre-2022	—	—	—	—	3.545	2.016	2.199
TabNet	Pre-2023	—	—	—	—	—	11.856	19.570
TabNet	Pre-2024	—	—	—	—	—	—	20.316
Anchor Regression (Causal)	Pre-2018	6.472	4.921	2.816	8.429	8.214	5.078	3.411
Anchor Regression (Causal)	Pre-2019	—	7.244	9.956	6.460	21.176	11.325	15.901
Anchor Regression (Causal)	Pre-2020	—	—	6.964	5.742	17.775	13.487	19.783
Anchor Regression (Causal)	Pre-2021	—	—	—	8.049	15.540	9.949	29.754
Anchor Regression (Causal)	Pre-2022	—	—	—	—	15.162	6.992	30.757
Anchor Regression (Causal)	Pre-2023	—	—	—	—	—	12.820	24.449
Anchor Regression (Causal)	Pre-2024	—	—	—	—	—	—	21.500

Table 19: **Max-of-Family dominance.**

Anchor Training	2018	2019	2020	2021	2022	2023	2024
Panel A: Maximum Absolute PR-AUC							
Pre-2018	0.986 (Causal)	0.734 (Linear)	0.604 (Deep)	0.473 (Causal)	0.582 (Trees)	0.597 (Trees)	0.667 (Linear)
Pre-2019	—	0.874 (Deep)	0.675 (Causal)	0.464 (Trees)	0.392 (Causal)	0.732 (Deep)	0.649 (Trees)
Pre-2020	—	—	0.642 (Trees)	0.519 (Deep)	0.424 (Trees)	0.655 (Deep)	0.399 (Linear)
Pre-2021	—	—	—	0.682 (Trees)	0.500 (Trees)	0.676 (Trees)	0.649 (Trees)
Pre-2022	—	—	—	—	0.607 (Trees)	0.689 (Trees)	0.572 (Causal)
Pre-2023	—	—	—	—	—	0.709 (Trees)	0.455 (Causal)
Pre-2024	—	—	—	—	—	—	0.421 (Trees)
Panel B: Maximum Relative PR-AUC Lift							
Pre-2018	6.47 (Causal)	9.07 (Linear)	8.91 (Deep)	8.43 (Causal)	31.44 (Trees)	17.14 (Trees)	35.83 (Linear)
Pre-2019	—	10.79 (Deep)	9.96 (Causal)	8.28 (Trees)	21.18 (Causal)	21.03 (Deep)	34.87 (Trees)
Pre-2020	—	—	9.47 (Trees)	9.25 (Deep)	22.91 (Trees)	18.81 (Deep)	21.44 (Linear)
Pre-2021	—	—	—	12.16 (Trees)	27.00 (Trees)	19.42 (Trees)	34.87 (Trees)
Pre-2022	—	—	—	—	32.79 (Trees)	19.79 (Trees)	30.76 (Causal)
Pre-2023	—	—	—	—	—	20.35 (Trees)	24.45 (Causal)
Pre-2024	—	—	—	—	—	—	22.65 (Trees)

B.2.4 Threshold: $>15\%$

Table 20: Temporal predictive drift: PR-AUC decay by algorithm.

Model	Anchor Training	2018	2019	2020	2021	2022	2023	2024
CatBoost	Pre-2018	0.937	0.433	0.642	0.399	0.592	0.718	0.649
CatBoost	Pre-2019	—	0.505	0.576	0.340	0.222	0.554	0.710
CatBoost	Pre-2020	—	—	0.656	0.331	0.262	0.525	0.395
CatBoost	Pre-2021	—	—	—	0.299	0.569	0.802	0.747
CatBoost	Pre-2022	—	—	—	—	0.508	0.663	0.401
CatBoost	Pre-2023	—	—	—	—	—	0.593	0.338
CatBoost	Pre-2024	—	—	—	—	—	—	0.401
XGBoost	Pre-2018	0.923	0.685	0.560	0.350	0.387	0.595	0.611
XGBoost	Pre-2019	—	0.467	0.660	0.325	0.331	0.481	0.604
XGBoost	Pre-2020	—	—	0.823	0.327	0.253	0.683	0.567
XGBoost	Pre-2021	—	—	—	0.290	0.285	0.431	0.560
XGBoost	Pre-2022	—	—	—	—	0.387	0.627	0.553
XGBoost	Pre-2023	—	—	—	—	—	0.436	0.404
XGBoost	Pre-2024	—	—	—	—	—	—	0.558
Logistic (L2)	Pre-2018	0.944	0.672	0.557	0.294	0.288	0.376	0.387
Logistic (L2)	Pre-2019	—	0.671	0.542	0.351	0.249	0.381	0.338
Logistic (L2)	Pre-2020	—	—	0.535	0.289	0.193	0.343	0.352
Logistic (L2)	Pre-2021	—	—	—	0.272	0.208	0.363	0.318
Logistic (L2)	Pre-2022	—	—	—	—	0.266	0.358	0.318
Logistic (L2)	Pre-2023	—	—	—	—	—	0.341	0.296
Logistic (L2)	Pre-2024	—	—	—	—	—	—	0.356
Spatial-FE Logistic	Pre-2018	0.714	0.681	0.507	0.287	0.283	0.479	0.622
Spatial-FE Logistic	Pre-2019	—	0.690	0.543	0.408	0.305	0.453	0.388
Spatial-FE Logistic	Pre-2020	—	—	0.552	0.296	0.194	0.452	0.349
Spatial-FE Logistic	Pre-2021	—	—	—	0.278	0.205	0.580	0.339
Spatial-FE Logistic	Pre-2022	—	—	—	—	0.261	0.429	0.327
Spatial-FE Logistic	Pre-2023	—	—	—	—	—	0.659	0.326
Spatial-FE Logistic	Pre-2024	—	—	—	—	—	—	0.401
TabNet	Pre-2018	0.885	0.636	0.646	0.237	0.261	0.307	0.458
TabNet	Pre-2019	—	0.524	0.652	0.394	0.200	0.167	0.179
TabNet	Pre-2020	—	—	0.679	0.389	0.258	0.348	0.371
TabNet	Pre-2021	—	—	—	0.350	0.525	0.342	0.754
TabNet	Pre-2022	—	—	—	—	0.667	0.392	0.854
TabNet	Pre-2023	—	—	—	—	—	0.255	0.507
TabNet	Pre-2024	—	—	—	—	—	—	0.399
Anchor Regression (Causal)	Pre-2018	0.912	0.560	0.499	0.286	0.306	0.494	0.432
Anchor Regression (Causal)	Pre-2019	—	0.627	0.603	0.308	0.230	0.372	0.450
Anchor Regression (Causal)	Pre-2020	—	—	0.402	0.318	0.293	0.258	0.499
Anchor Regression (Causal)	Pre-2021	—	—	—	0.328	0.278	0.277	0.536
Anchor Regression (Causal)	Pre-2022	—	—	—	—	0.250	0.428	0.500
Anchor Regression (Causal)	Pre-2023	—	—	—	—	—	0.182	0.405
Anchor Regression (Causal)	Pre-2024	—	—	—	—	—	—	0.365

Table 21: **Temporal predictive drift: PR-AUC lift by algorithm.**

Model	Anchor Training	2018	2019	2020	2021	2022	2023	2024
CatBoost	Pre-2018	6.151	6.493	9.474	8.530	42.600	24.055	34.874
CatBoost	Pre-2019	—	7.571	8.503	7.271	16.000	18.576	38.137
CatBoost	Pre-2020	—	—	9.679	7.078	18.857	17.592	21.212
CatBoost	Pre-2021	—	—	—	6.407	41.000	26.854	40.153
CatBoost	Pre-2022	—	—	—	—	36.545	22.202	21.569
CatBoost	Pre-2023	—	—	—	—	—	19.852	18.178
CatBoost	Pre-2024	—	—	—	—	—	—	21.569
XGBoost	Pre-2018	6.064	10.271	8.253	7.493	27.886	19.947	32.847
XGBoost	Pre-2019	—	7.012	9.740	6.947	23.857	16.122	32.474
XGBoost	Pre-2020	—	—	12.141	6.994	18.203	22.866	30.458
XGBoost	Pre-2021	—	—	—	6.203	20.545	14.434	30.074
XGBoost	Pre-2022	—	—	—	—	27.857	21.005	29.712
XGBoost	Pre-2023	—	—	—	—	—	14.617	21.692
XGBoost	Pre-2024	—	—	—	—	—	—	30.010
Logistic (L2)	Pre-2018	6.200	10.075	8.216	6.301	20.743	12.593	20.828
Logistic (L2)	Pre-2019	—	10.068	7.999	7.507	17.943	12.772	18.178
Logistic (L2)	Pre-2020	—	—	7.885	6.193	13.872	11.482	18.898
Logistic (L2)	Pre-2021	—	—	—	5.829	14.967	12.154	17.090
Logistic (L2)	Pre-2022	—	—	—	—	19.143	12.004	17.090
Logistic (L2)	Pre-2023	—	—	—	—	—	11.409	15.933
Logistic (L2)	Pre-2024	—	—	—	—	—	—	19.148
Spatial-FE Logistic	Pre-2018	4.692	10.219	7.473	6.145	20.396	16.030	33.434
Spatial-FE Logistic	Pre-2019	—	10.354	8.006	8.726	21.943	15.178	20.865
Spatial-FE Logistic	Pre-2020	—	—	8.146	6.328	13.943	15.149	18.733
Spatial-FE Logistic	Pre-2021	—	—	—	5.942	14.743	19.429	18.210
Spatial-FE Logistic	Pre-2022	—	—	—	—	18.800	14.387	17.581
Spatial-FE Logistic	Pre-2023	—	—	—	—	—	22.077	17.538
Spatial-FE Logistic	Pre-2024	—	—	—	—	—	—	21.569
TabNet	Pre-2018	5.812	9.542	9.535	5.072	18.800	10.295	24.635
TabNet	Pre-2019	—	7.858	9.612	8.430	14.369	5.596	9.630
TabNet	Pre-2020	—	—	10.019	8.317	18.545	11.659	19.943
TabNet	Pre-2021	—	—	—	7.483	37.835	11.461	40.536
TabNet	Pre-2022	—	—	—	—	48.000	13.148	45.911
TabNet	Pre-2023	—	—	—	—	—	8.551	27.259
TabNet	Pre-2024	—	—	—	—	—	—	21.425
Anchor Regression (Causal)	Pre-2018	5.986	8.407	7.355	6.111	22.000	16.535	23.228
Anchor Regression (Causal)	Pre-2019	—	9.403	8.895	6.589	16.533	12.448	24.188
Anchor Regression (Causal)	Pre-2020	—	—	5.934	6.795	21.091	8.645	26.795
Anchor Regression (Causal)	Pre-2021	—	—	—	7.019	20.000	9.267	28.816
Anchor Regression (Causal)	Pre-2022	—	—	—	—	18.000	14.331	26.875
Anchor Regression (Causal)	Pre-2023	—	—	—	—	—	6.083	21.750
Anchor Regression (Causal)	Pre-2024	—	—	—	—	—	—	19.623

Table 22: **Max-of-Family dominance.**

Anchor Training	2018	2019	2020	2021	2022	2023	2024
Panel A: Maximum Absolute PR-AUC							
Pre-2018	0.944 (Linear)	0.685 (Trees)	0.646 (Deep)	0.399 (Trees)	0.592 (Trees)	0.718 (Trees)	0.649 (Trees)
Pre-2019	—	0.690 (Linear)	0.660 (Trees)	0.408 (Linear)	0.331 (Trees)	0.554 (Trees)	0.710 (Trees)
Pre-2020	—	—	0.823 (Trees)	0.389 (Deep)	0.293 (Causal)	0.683 (Trees)	0.567 (Trees)
Pre-2021	—	—	—	0.350 (Deep)	0.569 (Trees)	0.802 (Trees)	0.754 (Deep)
Pre-2022	—	—	—	—	0.667 (Deep)	0.663 (Trees)	0.854 (Deep)
Pre-2023	—	—	—	—	—	0.659 (Linear)	0.507 (Deep)
Pre-2024	—	—	—	—	—	—	0.558 (Trees)
Panel B: Maximum Relative PR-AUC Lift							
Pre-2018	6.20 (Linear)	10.27 (Trees)	9.54 (Deep)	8.53 (Trees)	42.60 (Trees)	24.05 (Trees)	34.87 (Trees)
Pre-2019	—	10.35 (Linear)	9.74 (Trees)	8.73 (Linear)	23.86 (Trees)	18.58 (Trees)	38.14 (Trees)
Pre-2020	—	—	12.14 (Trees)	8.32 (Deep)	21.09 (Causal)	22.87 (Trees)	30.46 (Trees)
Pre-2021	—	—	—	7.48 (Deep)	41.00 (Trees)	26.85 (Trees)	40.54 (Deep)
Pre-2022	—	—	—	—	48.00 (Deep)	22.20 (Trees)	45.91 (Deep)
Pre-2023	—	—	—	—	—	22.08 (Linear)	27.26 (Deep)
Pre-2024	—	—	—	—	—	—	30.01 (Trees)

B.2.5 Threshold: $>20\%$

Table 23: Temporal predictive drift: PR-AUC decay by algorithm.

Model	Anchor Training	2018	2019	2020	2021	2022	2023	2024
CatBoost	Pre-2018	0.960	0.794	0.472	0.324	0.258	0.591	0.622
CatBoost	Pre-2019	—	0.524	0.481	0.405	0.143	0.771	0.395
CatBoost	Pre-2020	—	—	0.503	0.354	0.208	0.636	0.430
CatBoost	Pre-2021	—	—	—	0.467	0.325	0.799	0.415
CatBoost	Pre-2022	—	—	—	—	0.625	0.367	0.326
CatBoost	Pre-2023	—	—	—	—	—	0.633	0.443
CatBoost	Pre-2024	—	—	—	—	—	—	0.586
XGBoost	Pre-2018	0.926	0.553	0.467	0.341	0.208	0.485	0.579
XGBoost	Pre-2019	—	0.397	0.399	0.368	0.208	0.368	0.679
XGBoost	Pre-2020	—	—	0.511	0.276	0.196	0.543	0.458
XGBoost	Pre-2021	—	—	—	0.402	0.267	0.618	0.679
XGBoost	Pre-2022	—	—	—	—	0.250	0.443	0.679
XGBoost	Pre-2023	—	—	—	—	—	0.525	0.350
XGBoost	Pre-2024	—	—	—	—	—	—	0.442
Logistic (L2)	Pre-2018	0.852	0.700	0.413	0.236	0.139	0.277	0.771
Logistic (L2)	Pre-2019	—	0.667	0.358	0.280	0.132	0.244	0.360
Logistic (L2)	Pre-2020	—	—	0.347	0.281	0.133	0.302	0.367
Logistic (L2)	Pre-2021	—	—	—	0.264	0.162	0.258	0.278
Logistic (L2)	Pre-2022	—	—	—	—	0.202	0.348	0.296
Logistic (L2)	Pre-2023	—	—	—	—	—	0.360	0.320
Logistic (L2)	Pre-2024	—	—	—	—	—	—	0.305
Spatial-FE Logistic	Pre-2018	0.642	0.603	0.387	0.296	0.113	0.362	1.000
Spatial-FE Logistic	Pre-2019	—	0.694	0.403	0.341	0.127	0.416	0.374
Spatial-FE Logistic	Pre-2020	—	—	0.440	0.334	0.110	0.420	0.379
Spatial-FE Logistic	Pre-2021	—	—	—	0.285	0.134	0.304	0.305
Spatial-FE Logistic	Pre-2022	—	—	—	—	0.138	0.342	0.287
Spatial-FE Logistic	Pre-2023	—	—	—	—	—	0.471	0.298
Spatial-FE Logistic	Pre-2024	—	—	—	—	—	—	0.306
TabNet	Pre-2018	0.990	0.399	0.802	0.226	0.267	0.295	0.380
TabNet	Pre-2019	—	0.364	0.506	0.373	0.174	0.273	0.351
TabNet	Pre-2020	—	—	0.309	0.235	0.167	0.173	0.170
TabNet	Pre-2021	—	—	—	0.503	0.244	0.454	0.775
TabNet	Pre-2022	—	—	—	—	0.216	0.349	0.461
TabNet	Pre-2023	—	—	—	—	—	0.254	0.390
TabNet	Pre-2024	—	—	—	—	—	—	0.667
Anchor Regression (Causal)	Pre-2018	0.994	0.452	0.500	0.432	0.133	0.343	0.525
Anchor Regression (Causal)	Pre-2019	—	0.539	0.500	0.314	0.171	0.237	0.406
Anchor Regression (Causal)	Pre-2020	—	—	0.273	0.331	0.208	0.230	0.661
Anchor Regression (Causal)	Pre-2021	—	—	—	0.331	0.267	0.303	0.575
Anchor Regression (Causal)	Pre-2022	—	—	—	—	0.155	0.427	0.722
Anchor Regression (Causal)	Pre-2023	—	—	—	—	—	0.356	0.469
Anchor Regression (Causal)	Pre-2024	—	—	—	—	—	—	0.557

Table 24: **Temporal predictive drift: PR-AUC lift by algorithm.**

Model	Anchor Training	2018	2019	2020	2021	2022	2023	2024
CatBoost	Pre-2018	6.307	11.904	8.346	6.944	27.818	29.693	33.407
CatBoost	Pre-2019	—	7.861	8.512	8.671	15.429	38.734	21.212
CatBoost	Pre-2020	—	—	8.910	7.579	22.500	31.965	23.132
CatBoost	Pre-2021	—	—	—	10.000	35.100	40.130	22.332
CatBoost	Pre-2022	—	—	—	—	67.500	18.460	17.538
CatBoost	Pre-2023	—	—	—	—	—	31.825	23.825
CatBoost	Pre-2024	—	—	—	—	—	—	31.503
XGBoost	Pre-2018	6.084	8.301	8.267	7.302	22.500	24.384	31.130
XGBoost	Pre-2019	—	5.961	7.069	7.871	22.500	18.497	36.505
XGBoost	Pre-2020	—	—	9.053	5.909	21.214	27.279	24.635
XGBoost	Pre-2021	—	—	—	8.613	28.800	31.047	36.505
XGBoost	Pre-2022	—	—	—	—	27.000	22.270	36.505
XGBoost	Pre-2023	—	—	—	—	—	26.381	18.839
XGBoost	Pre-2024	—	—	—	—	—	—	23.740
Logistic (L2)	Pre-2018	5.595	10.494	7.317	5.040	15.058	13.895	41.432
Logistic (L2)	Pre-2019	—	10.001	6.345	5.997	14.308	12.240	19.330
Logistic (L2)	Pre-2020	—	—	6.149	6.003	14.400	15.172	19.708
Logistic (L2)	Pre-2021	—	—	—	5.648	17.532	12.945	14.964
Logistic (L2)	Pre-2022	—	—	—	—	21.808	17.475	15.901
Logistic (L2)	Pre-2023	—	—	—	—	—	18.082	17.204
Logistic (L2)	Pre-2024	—	—	—	—	—	—	16.381
Spatial-FE Logistic	Pre-2018	4.218	9.048	6.859	6.338	12.214	18.184	53.750
Spatial-FE Logistic	Pre-2019	—	10.412	7.139	7.306	13.708	20.905	20.116
Spatial-FE Logistic	Pre-2020	—	—	7.794	7.138	11.868	21.117	20.372
Spatial-FE Logistic	Pre-2021	—	—	—	6.092	14.464	15.282	16.381
Spatial-FE Logistic	Pre-2022	—	—	—	—	14.914	17.199	15.444
Spatial-FE Logistic	Pre-2023	—	—	—	—	—	23.689	16.004
Spatial-FE Logistic	Pre-2024	—	—	—	—	—	—	16.452
TabNet	Pre-2018	6.498	5.979	14.191	4.837	28.800	14.814	20.409
TabNet	Pre-2019	—	5.458	8.952	7.985	18.750	13.706	18.850
TabNet	Pre-2020	—	—	5.471	5.037	18.000	8.687	9.113
TabNet	Pre-2021	—	—	—	10.758	26.308	22.792	41.656
TabNet	Pre-2022	—	—	—	—	23.318	17.513	24.785
TabNet	Pre-2023	—	—	—	—	—	12.753	20.988
TabNet	Pre-2024	—	—	—	—	—	—	35.833
Anchor Regression (Causal)	Pre-2018	6.526	6.774	8.850	9.241	14.400	17.229	28.219
Anchor Regression (Causal)	Pre-2019	—	8.087	8.850	6.718	18.514	11.934	21.799
Anchor Regression (Causal)	Pre-2020	—	—	4.832	7.088	22.500	11.556	35.513
Anchor Regression (Causal)	Pre-2021	—	—	—	7.092	28.800	15.205	30.906
Anchor Regression (Causal)	Pre-2022	—	—	—	—	16.714	21.461	38.819
Anchor Regression (Causal)	Pre-2023	—	—	—	—	—	17.877	25.233
Anchor Regression (Causal)	Pre-2024	—	—	—	—	—	—	29.946

Table 25: **Max-of-Family dominance.**

Anchor Training	2018	2019	2020	2021	2022	2023	2024
Panel A: Maximum Absolute PR-AUC							
Pre-2018	0.994 (Causal)	0.794 (Trees)	0.802 (Deep)	0.432 (Causal)	0.267 (Deep)	0.591 (Trees)	1.000 (Linear)
Pre-2019	—	0.694 (Linear)	0.506 (Deep)	0.405 (Trees)	0.208 (Trees)	0.771 (Trees)	0.679 (Trees)
Pre-2020	—	—	0.511 (Trees)	0.354 (Trees)	0.208 (Causal)	0.636 (Trees)	0.661 (Causal)
Pre-2021	—	—	—	0.503 (Deep)	0.325 (Trees)	0.799 (Trees)	0.775 (Deep)
Pre-2022	—	—	—	—	0.625 (Trees)	0.443 (Trees)	0.722 (Causal)
Pre-2023	—	—	—	—	—	0.633 (Trees)	0.469 (Causal)
Pre-2024	—	—	—	—	—	—	0.667 (Deep)
Panel B: Maximum Relative PR-AUC Lift							
Pre-2018	6.53 (Causal)	11.90 (Trees)	14.19 (Deep)	9.24 (Causal)	28.80 (Deep)	29.69 (Trees)	53.75 (Linear)
Pre-2019	—	10.41 (Linear)	8.95 (Deep)	8.67 (Trees)	22.50 (Trees)	38.73 (Trees)	36.51 (Trees)
Pre-2020	—	—	9.05 (Trees)	7.58 (Trees)	22.50 (Causal)	31.96 (Trees)	35.51 (Causal)
Pre-2021	—	—	—	10.76 (Deep)	35.10 (Trees)	40.13 (Trees)	41.66 (Deep)
Pre-2022	—	—	—	—	67.50 (Trees)	22.27 (Trees)	38.82 (Causal)
Pre-2023	—	—	—	—	—	31.82 (Trees)	25.23 (Causal)
Pre-2024	—	—	—	—	—	—	35.83 (Deep)

B.2.6 Threshold: $>25\%$

Table 26: Temporal predictive drift: PR-AUC decay by algorithm.

Model	Anchor Training	2018	2019	2020	2021	2022	2023	2024
CatBoost	Pre-2018	0.869	0.476	0.574	0.371	0.111	0.388	0.950
CatBoost	Pre-2019	—	0.632	0.539	0.283	0.083	0.486	0.804
CatBoost	Pre-2020	—	—	0.414	0.266	0.125	0.494	0.747
CatBoost	Pre-2021	—	—	—	0.275	0.200	0.556	0.525
CatBoost	Pre-2022	—	—	—	—	0.500	0.645	0.590
CatBoost	Pre-2023	—	—	—	—	—	0.388	0.538
CatBoost	Pre-2024	—	—	—	—	—	—	0.799
XGBoost	Pre-2018	0.867	0.533	0.464	0.359	0.200	0.950	0.887
XGBoost	Pre-2019	—	0.490	0.361	0.389	0.091	0.360	0.950
XGBoost	Pre-2020	—	—	0.385	0.217	0.111	0.442	0.917
XGBoost	Pre-2021	—	—	—	0.357	0.143	0.476	0.679
XGBoost	Pre-2022	—	—	—	—	0.250	0.402	0.521
XGBoost	Pre-2023	—	—	—	—	—	0.318	0.617
XGBoost	Pre-2024	—	—	—	—	—	—	0.585
Logistic (L2)	Pre-2018	0.684	0.721	0.292	0.198	0.200	0.332	0.608
Logistic (L2)	Pre-2019	—	0.686	0.318	0.200	0.200	0.237	0.271
Logistic (L2)	Pre-2020	—	—	0.331	0.175	0.200	0.295	0.273
Logistic (L2)	Pre-2021	—	—	—	0.165	0.200	0.253	0.243
Logistic (L2)	Pre-2022	—	—	—	—	0.250	0.225	0.238
Logistic (L2)	Pre-2023	—	—	—	—	—	0.255	0.217
Logistic (L2)	Pre-2024	—	—	—	—	—	—	0.276
Spatial-FE Logistic	Pre-2018	0.475	0.621	0.333	0.169	0.167	0.335	0.679
Spatial-FE Logistic	Pre-2019	—	0.538	0.352	0.197	0.167	0.412	0.287
Spatial-FE Logistic	Pre-2020	—	—	0.383	0.148	0.100	0.412	0.229
Spatial-FE Logistic	Pre-2021	—	—	—	0.147	0.167	0.412	0.243
Spatial-FE Logistic	Pre-2022	—	—	—	—	0.250	0.295	0.228
Spatial-FE Logistic	Pre-2023	—	—	—	—	—	0.398	0.140
Spatial-FE Logistic	Pre-2024	—	—	—	—	—	—	0.178
TabNet	Pre-2018	0.889	0.707	0.365	0.182	0.100	0.245	0.729
TabNet	Pre-2019	—	0.623	0.313	0.182	0.100	0.124	0.500
TabNet	Pre-2020	—	—	0.587	0.183	0.005	0.028	0.056
TabNet	Pre-2021	—	—	—	0.406	0.083	0.552	0.608
TabNet	Pre-2022	—	—	—	—	0.053	0.067	0.131
TabNet	Pre-2023	—	—	—	—	—	0.402	0.469
TabNet	Pre-2024	—	—	—	—	—	—	0.830
Anchor Regression (Causal)	Pre-2018	0.797	0.068	0.126	0.321	0.250	0.205	0.255
Anchor Regression (Causal)	Pre-2019	—	0.303	0.422	0.226	0.077	0.267	0.350
Anchor Regression (Causal)	Pre-2020	—	—	0.248	0.233	0.077	0.430	0.583
Anchor Regression (Causal)	Pre-2021	—	—	—	0.288	0.091	0.251	0.461
Anchor Regression (Causal)	Pre-2022	—	—	—	—	0.250	0.222	0.406
Anchor Regression (Causal)	Pre-2023	—	—	—	—	—	0.279	0.472
Anchor Regression (Causal)	Pre-2024	—	—	—	—	—	—	0.384

Table 27: Temporal predictive drift: PR-AUC lift by algorithm.

Model	Anchor Training	2018	2019	2020	2021	2022	2023	2024
CatBoost	Pre-2018	6.111	7.692	11.287	13.243	24.000	19.472	51.062
CatBoost	Pre-2019	—	10.211	10.604	10.096	18.000	24.427	43.224
CatBoost	Pre-2020	—	—	8.141	9.475	27.000	24.803	40.153
CatBoost	Pre-2021	—	—	—	9.824	43.200	27.944	28.219
CatBoost	Pre-2022	—	—	—	—	108.000	32.417	31.727
CatBoost	Pre-2023	—	—	—	—	—	19.472	28.928
CatBoost	Pre-2024	—	—	—	—	—	—	42.925
XGBoost	Pre-2018	6.099	8.617	9.120	12.808	43.200	47.737	47.703
XGBoost	Pre-2019	—	7.913	7.108	13.876	19.636	18.071	51.062
XGBoost	Pre-2020	—	—	7.570	7.741	24.000	22.194	49.271
XGBoost	Pre-2021	—	—	—	12.747	30.857	23.929	36.505
XGBoost	Pre-2022	—	—	—	—	54.000	20.180	27.995
XGBoost	Pre-2023	—	—	—	—	—	15.977	33.172
XGBoost	Pre-2024	—	—	—	—	—	—	31.418
Logistic (L2)	Pre-2018	4.810	11.645	5.739	7.050	43.200	16.707	32.698
Logistic (L2)	Pre-2019	—	11.075	6.250	7.147	43.200	11.913	14.589
Logistic (L2)	Pre-2020	—	—	6.519	6.235	43.200	14.814	14.671
Logistic (L2)	Pre-2021	—	—	—	5.884	43.200	12.721	13.078
Logistic (L2)	Pre-2022	—	—	—	—	54.000	11.287	12.810
Logistic (L2)	Pre-2023	—	—	—	—	—	12.812	11.646
Logistic (L2)	Pre-2024	—	—	—	—	—	—	14.840
Spatial-FE Logistic	Pre-2018	3.343	10.031	6.545	6.034	36.000	16.842	36.505
Spatial-FE Logistic	Pre-2019	—	8.691	6.930	7.030	36.000	20.696	15.444
Spatial-FE Logistic	Pre-2020	—	—	7.530	5.296	21.600	20.696	12.318
Spatial-FE Logistic	Pre-2021	—	—	—	5.227	36.000	20.698	13.050
Spatial-FE Logistic	Pre-2022	—	—	—	—	54.000	14.836	12.280
Spatial-FE Logistic	Pre-2023	—	—	—	—	—	20.010	7.533
Spatial-FE Logistic	Pre-2024	—	—	—	—	—	—	9.564
TabNet	Pre-2018	6.256	11.421	7.171	6.499	21.600	12.323	39.193
TabNet	Pre-2019	—	10.056	6.152	6.503	21.600	6.236	26.875
TabNet	Pre-2020	—	—	11.535	6.519	1.000	1.411	2.990
TabNet	Pre-2021	—	—	—	14.474	18.000	27.717	32.698
TabNet	Pre-2022	—	—	—	—	11.368	3.374	7.060
TabNet	Pre-2023	—	—	—	—	—	20.205	25.233
TabNet	Pre-2024	—	—	—	—	—	—	44.632
Anchor Regression (Causal)	Pre-2018	5.610	1.101	2.488	11.465	54.000	10.300	13.688
Anchor Regression (Causal)	Pre-2019	—	4.896	8.304	8.067	16.615	13.400	18.812
Anchor Regression (Causal)	Pre-2020	—	—	4.875	8.316	16.615	21.617	31.354
Anchor Regression (Causal)	Pre-2021	—	—	—	10.274	19.636	12.599	24.785
Anchor Regression (Causal)	Pre-2022	—	—	—	—	54.000	11.167	21.831
Anchor Regression (Causal)	Pre-2023	—	—	—	—	—	14.028	25.382
Anchor Regression (Causal)	Pre-2024	—	—	—	—	—	—	20.656

Table 28: Max-of-Family dominance.

Anchor Training	2018	2019	2020	2021	2022	2023	2024
Panel A: Maximum Absolute PR-AUC							
Pre-2018	0.889 (Deep)	0.721 (Linear)	0.574 (Trees)	0.371 (Trees)	0.250 (Causal)	0.950 (Trees)	0.950 (Trees)
Pre-2019	—	0.686 (Linear)	0.539 (Trees)	0.389 (Trees)	0.200 (Linear)	0.486 (Trees)	0.950 (Trees)
Pre-2020	—	—	0.587 (Deep)	0.266 (Trees)	0.200 (Linear)	0.494 (Trees)	0.917 (Trees)
Pre-2021	—	—	—	0.406 (Deep)	0.200 (Linear)	0.556 (Trees)	0.679 (Trees)
Pre-2022	—	—	—	—	0.500 (Trees)	0.645 (Trees)	0.590 (Trees)
Pre-2023	—	—	—	—	—	0.402 (Deep)	0.617 (Trees)
Pre-2024	—	—	—	—	—	—	0.830 (Deep)
Panel B: Maximum Relative PR-AUC Lift							
Pre-2018	6.26 (Deep)	11.65 (Linear)	11.29 (Trees)	13.24 (Trees)	54.00 (Causal)	47.74 (Trees)	51.06 (Trees)
Pre-2019	—	11.07 (Linear)	10.60 (Trees)	13.88 (Trees)	43.20 (Linear)	24.43 (Trees)	51.06 (Trees)
Pre-2020	—	—	11.54 (Deep)	9.47 (Trees)	43.20 (Linear)	24.80 (Trees)	49.27 (Trees)
Pre-2021	—	—	—	14.47 (Deep)	43.20 (Linear)	27.94 (Trees)	36.51 (Trees)
Pre-2022	—	—	—	—	108.00 (Trees)	32.42 (Trees)	31.73 (Trees)
Pre-2023	—	—	—	—	—	20.20 (Deep)	33.17 (Trees)
Pre-2024	—	—	—	—	—	—	44.63 (Deep)

B.3 SHAP Beeswarm Exhibit

This appendix reports the case-level SHAP beeswarm for the primary filing-date model as a descriptive model-reliance exhibit only. It is included to show the within-model distribution of signed feature contributions, not to support causal interpretation or a portable substantive narrative.




Figure 16: SHAP beeswarm for the filing-date opposition model. Each point is a case-level contribution for one predictor; horizontal position gives the signed effect on the model score and vertical density shows the distribution of contributions across cases. Because correlated predictors can exchange attribution weight, this exhibit is descriptive support for within-model behavior rather than a stand-alone substantive ranking of stable feature importance.

C Institutional Context and Causal Overlays

C.1 Institutional Geographic Overlay Analyses

Additional analyses evaluate localized institutional overlays, Neighborhood Plan Areas (NPA), Historic District (HD) designations, Transit-Oriented Development (TOD) deregulation, and the 10-1 council transition system.

Threshold-Based Reduced-Form Discontinuity on Reconstructed Petition Share




Figure 17: Supplementary identification schematic for the threshold-based reduced-form discontinuity design. The running variable R (reconstructed signed-area share) defines threshold crossing T ($\geq 20\%$), and the model estimates the associated discontinuous shift in council processing delay Y without observing clerk certification compliance. Unobserved factors U may influence both petition mobilization and delay through channels not mediated by the threshold.

The ITS check for the 10-1 transition reports coefficient 0.041 on council dissent ($p = 0.74$), and the NPA check reports coefficient 0.083 ($p = 0.61$). Historic District and TOD overlay designs are not estimable in this panel because outcome and/or treatment variation collapses once the legally relevant window restrictions are imposed. These outputs are retained only for transparency and are not interpreted substantively.

2022 Electoral Transition: Geographic Treatment Effects




Figure 18: Geographic choropleth illustrating the descriptive interaction estimate. Districts 4 and 9 map the estimated coefficient ($\Delta - 0.546$), showing the interaction effect for districts that changed representation, relative to districts without comparable turnover.

Placebo Falsification: Petition Filing Rate by Election Cycle




Figure 19: Placebo estimates across prior municipal electoral cycles (2014, 2016, 2018, 2020) for the affected districts. Only the 2022 cycle shows a detectable effect.

Table 29: Summary of Supplementary Quasi-Experimental Estimates for Reconstructed Petition Outcomes

Design	Dep. Var.	Coeff.		SE	<i>p</i>	<i>N</i>
(1) 10-1 Council ITS	vote_no	0.041		0.122	0.74	20
(2) NPA Designation	vote_no	0.083		0.245	0.61	20
(3) Historic District	vote_no	not estimable	not estimable	not estimable	not estimable	20
(4) TOD Deregulation	vote_no	not estimable	not estimable	not estimable	not estimable	20
(5) 2022 Flipped District DiD	measured_threshold_crossing	+0.498		0.266	0.0136	518

OLS with robust standard errors where estimable. Designs (3) and (4) are marked "not estimable" due to collapsed identifying variation (zero/near-zero treatment or outcome support) in the restricted institutional windows. Design (5) reports the Flipped $\times$ later-window interaction coefficient.

D Policy Event Studies

D.1 HOME Initiative Event Study

Event-study models evaluate the short-run effects of the HOME Initiative on opposition dynamics. HOME Phase 1 and Phase 2 are treated as independent policy shocks, with treatment defined at the case level by parcel eligibility and timing anchored to application-acceptance start dates [14, 8].

Given the limited post-policy window and missing defensible controls, this sample does not support an estimable causal HOME effect. For transparency, descriptive TWFE coefficients are retained in the replication outputs, but they are not interpreted as policy effects in light of the well-known heterogeneous-timing problems in TWFE designs [26, 52]. In this manuscript, HOME results are explicitly marked "not estimable" for identification purposes, with figure-based event timing shown only as context [11].

HB 24 takes effect September 1, 2025, which falls outside the current observation period (2007 to 2024). Its evaluation is explicitly pre-registered as a future extension.




Figure 20: HOME Phase 1 event-study coefficients for measured threshold-crossing petitions relative to the implementation date. The absence of a clean spatial control cohort necessitates purely temporal event-tracking.